

Integration Methods in cppDE

Simon Beyer

2026-09-15

Contents

1	Overview	2
2	cppODE()	2
2.1	Initial value problem and four methods	2
2.2	Stiffness and the choice of method	3
2.3	Multistep methods	3
2.4	Single-step methods	4
2.5	Adaptive step-size control	6
2.6	Dense output	7
2.7	Automatic Forward sensitivities	7
2.8	Events and restarts	9
2.9	Reverse-mode sensitivities: the discrete adjoint	13
2.10	Method selection	22
3	cvode()	22
3.1	What is shared	23
3.2	What differs	23
3.3	Adjoint sensitivity analysis	23
4	cppFUN()	24
4.1	Unified API and two computational paths	24
4.2	Chain rule via dX, dP, dX2, dP2	25
4.3	Pass-through of unmodelled inputs	26
5	Benchmarks	27
	References	29
A	BDF and NDF coefficients	30
B	Adams-Moulton coefficients	30
C	Rosenbrock4 coefficients	30
C.1	Diagonal coefficient, time partials, and node positions	31
C.2	Stage couplings α_{ij}	31
C.3	Stage couplings γ_{ij}	31
C.4	Solution and error weights	31
C.5	Dense-output coefficients	32
D	Tsitouras 5(4) coefficients	32

1 Overview

cppDE generates and compiles per-model C++ code for three related problems through three entry points that share a common R interface: `cppODE()`, `cvode()`, and `cppFUN()`. Table 1 summarises their scope.

	<code>cppODE()</code>	<code>cvode()</code>	<code>cppFUN()</code>
Problem class	ODE	ODE	algebraic $y = g(x, p)$
Methods	<code>bdf</code> , <code>adams</code> , <code>rb4</code> , <code>tsit5</code>	<code>bdf</code> , <code>adams</code>	
First-order sensitivities	in-tree dual-number AD	SUNDIALS analytic RHS	dual AD or symbolic
Reverse-mode gradient	yes (discrete adjoint)	yes (CVODES ASA)	yes
Second-order sensitivities	yes (nested duals, fwd-over-rev)	no	yes (nested duals or symbolic)
External dependencies	LAPACK; KLU optional	SUNDIALS ≥ 6.0 ; KLU opt.	none
Codegen	SymPy + R CMD SHLIB	SymPy + R CMD SHLIB	SymPy + R CMD SHLIB

Table 1: The three cppDE entry points and what they cover. The codegen pipeline (parse the user’s expressions through SymPy, emit C++, compile via R CMD SHLIB, load via `dyn.load`) is shared, as is the `solveODE()` dispatcher used by both integrator backends.

The integrator entry points `cppODE()` and `cvode()` expose the same R-side interface: the user supplies named character vectors of right-hand sides together with optional events, forcings, and parameter vectors, and the returned model handle is passed to `solveODE()`. The two backends differ only in the numerical core. `cppODE()` emits a self-contained C++ source that includes the headers under `inst/include/cppde/`; the resulting binary depends on LAPACK and, optionally, KLU. `cvode()` emits a thin glue layer that links against the system-installed SUNDIALS CVODE/CVODES library and inherits its numerical behaviour.

Section 2 describes the in-tree integrator `cppODE()` in full, including its four time-stepping methods, its forward-mode automatic-differentiation framework for parameter sensitivities, the reverse-mode discrete adjoint that computes a gradient at a cost independent of the number of parameters, and the transport of sensitivities across events. Section 3 describes the SUNDIALS-backed `cvode()` integrator and the adjoint sensitivity analysis it reaches the same gradient by. Section 4 describes the algebraic-function compiler `cppFUN()`, which compiles a system of expressions $y = g(x, p)$ together with optional first- and second-order derivatives, with no time integration. The principal use cases of `cppFUN()` are observation maps that link integrator output to data in likelihood-based inference and reparametrisation Jacobians $\partial p / \partial \theta$ that seed sensitivity propagation through the integrators. Section 5 benchmarks the two integrator backends against each other on 42 problems.

2 cppODE()

2.1 Initial value problem and four methods

This section describes the in-tree integrator of `cppODE()` in detail: the four time-stepping methods, the adaptive step-size control they share, dense Hermite output, sparse and dense LU factorisation, time- and root-triggered events, and forward-mode sensitivity propagation through dual numbers. The methods differ only in how the individual time step is taken. The Butcher tableaus and numerical coefficients of each method are reproduced in Appendices A–D.

The initial-value problem solved by `cppODE()` is

$$\dot{x} = f(x, p_{\text{dyn}}, t), \quad x(t_0) = p_{\text{init}}, \quad (1)$$

with state $x \in \mathbb{R}^{n_x}$ and parameter vector $p = (p_{\text{init}}, p_{\text{dyn}})$ collecting the initial conditions and the dynamic parameters. The right-hand side f may depend explicitly on time, in which case the system is non-autonomous; an explicit time argument is also necessary for the forcing-function machinery used to inject time-dependent inputs. The Jacobian of f with respect to the state is denoted $J = \partial f / \partial x$ throughout.

2.2 Stiffness and the choice of method

The choice between the four methods is governed primarily by whether the problem is *stiff*. A problem is stiff on a given integration interval when its solution contains components evolving on time scales much shorter than the interval length. An explicit integration method applied to such a problem is forced to take step sizes set by the fastest component rather than by the accuracy requirement, leading to step counts that are uneconomic in practice. Implicit methods avoid this constraint by absorbing the fast components into a linear solve at each step.

The standard tool for analysing the stability of a method is the linear scalar test equation $\dot{y} = \lambda y$ with $\lambda \in \mathbb{C}^-$, whose exact solution decays. Substituting this equation into a one-step method produces a single scalar amplification factor $R(z)$ acting on y per step, where $z = h\lambda$ combines the step size and the eigenvalue. The method is *A-stable* if $|R(z)| \leq 1$ on the entire left half-plane, so that the numerical solution decays whenever the exact one does, regardless of step size; *A(α)-stable* if the same bound holds on the wedge $\{z \in \mathbb{C} : |\arg(-z)| \leq \alpha\}$ for some opening angle $\alpha < 90^\circ$; and *L-stable* if it is A-stable and additionally $R(z) \rightarrow 0$ as $\text{Re}(z) \rightarrow -\infty$, the latter condition being what damps very stiff transients out in a single step rather than carrying them along.

The initial-value problem (1) is called *stiff* on $[t_0, t_{\text{end}}]$ when an explicit method is forced to take step sizes much smaller than what local accuracy alone would demand, in order to keep the dimensionless products $h\lambda_k$ inside the stability domain, where $\{\lambda_k\}$ are the eigenvalues of J along the trajectory. Stiffness is therefore not a property of the system in isolation; it depends on the spectrum of J , on the integration horizon, and on the requested tolerance. Chemical-kinetics networks with rate constants spanning several orders of magnitude are stiff in this sense, as are many discretised parabolic PDEs.

Stiff problems require methods whose stability domain contains the entire left half-plane (or a wide wedge thereof). Of the four methods listed above, BDF/NDF and Rosenbrock4 satisfy this requirement; Adams-Moulton and Tsit5 do not, and will perform poorly on stiff systems regardless of the requested tolerance.

2.3 Multistep methods

Linear multistep methods build the solution at the new time step x_n from a linear combination of the most recent $k + 1$ states and right-hand-side values, $\{x_{n-j}, f_{n-j}\}_{j=0}^k$. They re-use information from earlier steps and so attain high accuracy with few right-hand-side evaluations per step, at the price of having to manage that history.

cppDE implements both the BDF and the Adams families through a common internal representation, due to Nordsieck and used throughout SUNDIALS [1]. The idea is to store, instead of the past $k + 1$ states, a single *scaled derivative array*

$$z_n = \left(x_n, h_n \dot{x}_n, \frac{h_n^2}{2!} \ddot{x}_n, \dots, \frac{h_n^k}{k!} x_n^{(k)} \right),$$

i.e. the value at the current time and its first k derivatives, each multiplied by the appropriate power of the current step size to make the entries dimensionally homogeneous. Changing the step size or the integration order then becomes a single linear transformation of z_n rather than a refit through the past data, which simplifies both the bookkeeping and the implementation of variable-order and variable-step control.

Backward Differentiation Formula (bdf)

The BDF formula of order q defines x_n implicitly by the condition

$$\sum_{j=0}^q \alpha_j x_{n-j} = h_n f(x_n, t_n), \quad (2)$$

which is exact whenever $x(t)$ is a polynomial of degree at most q on the past $q + 1$ nodes. The method is A-stable for $q \leq 2$ and A(α)-stable for $q \in \{3, 4, 5\}$; the stability angle α shrinks monotonically from 86° at $q = 3$ to 51° at $q = 5$.

The implicit corrector (2) is solved by Newton iteration,

$$W \Delta x^{(\nu)} = -\tilde{r}(x_n^{(\nu)}), \quad x_n^{(\nu+1)} = x_n^{(\nu)} + \Delta x^{(\nu)}, \quad (3)$$

with leading coefficient $\beta_0 = 1/\alpha_0$ and the iteration matrix

$$W = \frac{1}{h_n \beta_0} \mathbb{1} - J(x_n^{(\nu)}, t_n), \quad (4)$$

in the form adopted by SUNDIALS [1], which differs from the textbook expression $I - h_n \beta_0 J$ by the overall scalar factor $h_n \beta_0$. The factorisation of W is the dominant cost per accepted step and is reused for as long as the contraction rate of the Newton iteration permits.

The default in cppDE is the Numerical Differentiation Formula (NDF) variant of [2], which adds a correction proportional to the Nordsieck predictor error to the BDF corrector,

$$\sum_{m=1}^q \frac{1}{m} \nabla^m x_n - h_n f(x_n, t_n) - \kappa_q \gamma_q (x_n - x_n^{\text{pred}}) = 0, \quad \gamma_q = \sum_{m=1}^q \frac{1}{m}, \quad (5)$$

where ∇ denotes the backward-difference operator. The algebraic effect of the modification is to scale β_0 in (3) by a factor $(1 - \kappa_q)$ and to multiply the leading truncation-error coefficient $1/(q+1)$ by the factor $1 + \kappa_q(q+1)\gamma_q$. With $\kappa_q < 0$ the truncation error decreases at fixed step size while the iteration matrix is unchanged, which permits a larger step at a fixed accuracy requirement. The values of κ_q adopted by cppDE are tabulated in Appendix A; they reproduce the 25% step-size gain over classical BDF reported by [2] for $q \leq 4$, at the cost of a small reduction in stability angle. Setting `useNDF = FALSE` recovers the classical BDF formulas ($\kappa_q \equiv 0$).

A typical accepted step performs one right-hand-side evaluation for the predictor and between one and three right-hand-side evaluations together with the same number of linear solves for the corrector. A new Jacobian evaluation and LU factorisation are triggered only when the Newton iteration stalls or when the step size or order changes substantially. BDF/NDF is the default method in cppDE and the recommended choice when stiffness is suspected.

Adams-Moulton in PECE form (adams)

The Adams-Moulton family integrates the local interpolant of f over $[t_{n-1}, t_n]$. Writing f as the unique polynomial through the $q+1$ samples $\{(t_{n-j}, f_{n-j})\}_{j=0}^q$ and integrating gives

$$x_n = x_{n-1} + h_n \sum_{j=0}^q \beta_j f(x_{n-j}, t_{n-j}), \quad (6)$$

which is implicit in x_n through the $j=0$ term. cppDE evaluates the corrector in predict-evaluate-correct-evaluate (PECE) mode rather than by Newton iteration. An explicit predictor (Adams-Bashforth of the same order q) first supplies a candidate value x_n^P for the new state; the right-hand side is then evaluated at x_n^P ; the corrector (6) is evaluated explicitly with $f(x_n^P, t_n)$ in place of $f(x_n, t_n)$, yielding a refined x_n ; and a final right-hand-side evaluation at the refined state is entered into the Nordsieck history. The procedure costs two right-hand-side evaluations per accepted step, with no nonlinear iteration and no linear solve.

The Adams-Moulton stability domain on the negative real axis shrinks rapidly with order; for $q \geq 3$ it does not contain even a small wedge of $\mathbb{C}^-$ away from the origin, so the method is unsuitable for stiff problems. For non-stiff problems with smooth solutions and a long integration horizon, however, the high attainable order ($q_{\text{max}} = 12$) and the cost of two right-hand-side evaluations per accepted step make it competitive with explicit Runge-Kutta methods of fixed order.

2.4 Single-step methods

Single-step methods carry no information beyond the current state. They achieve order $p > 1$ by evaluating internal stages within each step rather than by reaching back into the integration history. The absence of history makes restarts after events trivial; both methods below restart at full order in a single step, unlike the multistep methods, which must rebuild the Nordsieck history at order 1.

Rosenbrock (rb4)

A Rosenbrock method of s stages with diagonal coefficient γ linearises the implicit Runge-Kutta condition about (x_n, t_n) and defines stage values k_i by

$$Wk_i = f \left(x_n + \sum_{j<i} \alpha_{ij} k_j, t_n + c_i h_n \right) + \frac{1}{h_n} \sum_{j<i} \frac{\gamma_{ij}}{\gamma} k_j + h_n d_i \frac{\partial f}{\partial t}(x_n, t_n), \quad (7)$$

with the iteration matrix

$$W = \frac{1}{\gamma h_n} \mathbb{1} - J(x_n, t_n). \quad (8)$$

The method coefficients are listed in Table 2.

Symbol	Meaning
α_{ij}	weights with which stage i combines all earlier stages k_j , $j < i$
c_i	time abscissa of stage i inside the step, with $c_i = \sum_{j<i} \alpha_{ij}$
γ_{ij}	weights with which the linear-solve correction at stage i uses earlier stages
γ	common diagonal coefficient appearing in the iteration matrix W
d_i	weights that pick up the explicit time dependence of f via $\partial f / \partial t$
m_i	weights combining the stages into the advanced solution x_{n+1}
$\hat{m}_i$	weights combining the stages into the embedded order-3 estimator used for step control

Table 2: Rosenbrock method coefficients.

The advanced solution and embedded error estimate are

$$x_{n+1} = x_n + \sum_{i=1}^s m_i k_i, \quad \hat{x}_{n+1} = x_n + \sum_{i=1}^s \hat{m}_i k_i.$$

Each stage (7) is a single linear solve; no Newton iteration is required. Per accepted step `cppDE` evaluates J and factorises W once, performs five right-hand-side evaluations, and back-substitutes six times.

`cppDE` uses the order-4 coefficients of [3, Sec. IV.7], the same set as the `rosenbrock4` stepper of `Boost.NumericOdeint`. The pair is L-stable and carries an embedded order-3 estimator. Continuous output on $[t_n, t_{n+1}]$ is delivered by an order-3 interpolant constructed from the stage values k_i and two additional coefficient sets $d_i^{(2)}$ and $d_i^{(3)}$ tabulated in Appendix C; no extra right-hand-side evaluations are required for output points.

Rosenbrock methods are competitive with BDF on stiff problems whose Jacobian is cheap relative to the right-hand side, and they avoid the order-rebuilding cost incurred by the multistep methods after an event. Rosenbrock4 is therefore the recommended stiff method for systems with frequent state discontinuities at moderate state dimension.

Tsitouras (tsit5)

Tsitouras is the explicit Runge-Kutta pair of order 5(4) derived by [4] under the simplifying assumption that the first column of the coefficient matrix vanishes, $a_{i,1} = 0$ for $i = 2, \dots, s$. Removing this column reduces the number of order conditions to be satisfied during construction and was shown to admit a particularly favourable choice of nodes c_i in terms of leading-error norm. The pair has $s = 7$ stages with the First-Same-As-Last (FSAL) property,

$$k_7 = f(x_{n+1}, t_n + h_n) = k_1 \text{ at step } n + 1,$$

so an accepted step costs six right-hand-side evaluations rather than seven. The method coefficients are listed in Table 3.

Symbol	Meaning
a_{ij}	weights with which stage i combines earlier stages (lower triangular; first column is zero)
c_i	time abscissa of stage i inside the step, with $c_i = \sum_{j<i} a_{ij}$
b_i	weights combining the stages into the advanced order-5 solution; $b_7 = 0$
$\hat{b}_i$	weights combining the stages into the embedded order-4 estimator; $\hat{b}_7 \neq 0$
k_i	right-hand side at stage i : $k_i = f\left(x_n + h_n \sum_{j<i} a_{ij} k_j, t_n + c_i h_n\right)$

Table 3: Tsit5 method coefficients.

The advanced solution and embedded error estimate are

$$x_{n+1} = x_n + h_n \sum_{i=1}^7 b_i k_i, \quad \hat{x}_{n+1} = x_n + h_n \sum_{i=1}^7 \hat{b}_i k_i,$$

and the local error estimate used for step-size control is

$$e_{n+1} = h_n \sum_{i=1}^7 (\hat{b}_i - b_i) k_i.$$

Continuous output uses the accepted-step endpoints and their derivatives $(f_n, f_{n+1}) = (k_1, k_7)$, both already in cache from the FSAL property, so an output point costs no further right-hand-side evaluation.

Tsitouras has no Jacobian and no linear solve and performs the largest amount of integration per right-hand-side call of the four methods considered here. It is the appropriate choice for non-stiff problems whose right-hand side is expensive to evaluate.

2.5 Adaptive step-size control

The error estimate e_n delivered by each method is reduced to a scalar through the weighted root-mean-square norm

$$\|e_n\|_{\text{WRMS}} = \sqrt{\frac{1}{n_x} \sum_{j=1}^{n_x} \left(\frac{e_{n,j}}{w_{n,j}} \right)^2}, \quad w_{n,j} = \text{atol} + \text{rtol} |x_{n,j}^{\text{ref}}|, \quad (9)$$

in which $w_{n,j}$ is a per-component tolerance combining an absolute floor **atol** and a relative scale **rtol** $|x_{n,j}^{\text{ref}}|$. Normalising $e_{n,j}$ by $w_{n,j}$ before squaring places components of widely differing magnitudes on a common scale and removes the dependence of the acceptance test on the physical units of the state. The reference state x_n^{ref} that enters the relative term is the Nordsieck predictor in BDF/NDF and the component-wise maximum $\max(|x_n|, |x_{n+1}|)$ in the single-step methods; the two choices are both standard and produce numerically indistinguishable values of the norm in practice. A step is accepted when $\|e_n\|_{\text{WRMS}} \leq 1$.

The next step size h_{n+1} is selected by a control law of the type analysed in [3], applied to the WRMS norm of the current step and, when active, the WRMS norm of the previous step,

$$h_{n+1} = h_n \min \left\{ \rho_{\max}, \max \left(\rho_{\min}, S \|e_n\|_{\text{WRMS}}^{-1/(p+1)} \|e_{n-1}\|_{\text{WRMS}}^{\beta/(p+1)} \right) \right\},$$

with safety factor $S < 1$, growth limits $\rho_{\min}, \rho_{\max}$, and a memory coefficient $\beta \in [0, 1]$. The case $\beta = 0$ is the standard integral or I-controller, in which h_{n+1} depends on the current error norm only; cppDE uses this form for the multistep methods, where the Nordsieck representation already smooths the underlying signal. The case $\beta > 0$ is the proportional-integral (PI) controller of [3]; the single-step methods adopt this form to damp the high-frequency oscillations in h_n that a pure I-controller can otherwise produce on stiff or near-singular problems.

The first step

The controller needs a step to correct, and the first one has no error norm behind it. cppDE estimates it from the initial state in the manner of [5], Algorithm II.4.4: a rough $h_0 \sim \|x_0\|/\|f_0\|$ fixes a scale, a second derivative $\ddot{x}$ along the trajectory fixes the curvature, and

$$h_1 = \left(\frac{0.01}{\max(\|f_0\|, \|\ddot{x}\|)} \right)^{1/(p+1)} \quad (10)$$

turns the two into a step for a method of order p . The implicit methods obtain $\ddot{x} = \partial f / \partial t + Jf$ from the generated Jacobian, the explicit one from a finite difference of f . All norms are the weighted norms of (9), so for an AD value type the derivative components enter as well and the sensitivity dynamics influence the first step.

Two bounds follow, in the form CVODE uses. An upper bound keeps a single Euler step from changing any component by more than a fraction of its own scale, and is capped by the remaining interval; a lower bound keeps h_0 above the rounding of the time variable itself. The multistep methods reach the same estimate through a port of the CVODES routine `cvHin`, which refines $\sqrt{h_{\text{lb}} h_{\text{ub}}}$ by a finite-difference curvature until two successive estimates agree within a factor of two, so that the first step cppDE takes is the one CVODES would take on the same problem. The same estimator runs again after every event, on the post-event state.

2.6 Dense output

User-requested output times do not in general coincide with the adaptive step grid produced by the controller. Each accepted step therefore carries a local interpolant that allows evaluation of $x(t)$ at any time inside that step at the cost of polynomial arithmetic only. cppDE uses Hermite interpolation, which constrains both values and derivatives of the interpolant at the step endpoints. This form attains the same order as the underlying step without recourse to data outside the step: state and right-hand side at the two endpoints, augmented by stage values for the single-step methods. No additional right-hand-side evaluations are required for dense output. The local interpolation order is 4 for Rosenbrock4 and Tsit5; the multistep methods reuse the Nordsieck history, which carries derivatives up to order q .

2.7 Automatic Forward sensitivities

Dual-number arithmetic

Parameter sensitivities of the trajectory are computed by forward-mode automatic differentiation rather than by finite differences (whose truncation and round-off errors require multiple integrations to control) or by symbolic derivation of the sensitivity equations (whose intermediate expressions grow rapidly with the dimension of f). Forward-mode AD augments every floating-point scalar in the computation with an additional component that records its directional derivative with respect to a chosen parameter direction. The arithmetic operations and elementary functions are overloaded to act on this augmented type so that the chain rule is applied automatically; the C++ source that evaluates f is then compiled unchanged for both the primal and the sensitivity computation.

The augmented type is the *dual number*

$$\tilde{x} = x + x' \varepsilon, \quad \varepsilon^2 = 0, \quad \varepsilon \neq 0,$$

where $x \in \mathbb{R}$ is the primal value and $x' \in \mathbb{R}$ is the directional derivative. Treating $\tilde{x}$ as the first two terms of a Taylor expansion in ε and discarding all higher powers (which vanish by $\varepsilon^2 = 0$), the arithmetic of dual numbers follows from the standard rules of calculus: for $\tilde{x} = x + x' \varepsilon$ and $\tilde{y} = y + y' \varepsilon$,

$$\tilde{x} + \tilde{y} = (x + y) + (x' + y') \varepsilon, \quad \tilde{x} \cdot \tilde{y} = xy + (xy' + x'y) \varepsilon,$$

and for any smooth scalar function g ,

$$g(\tilde{x}) = g(x) + g'(x) x' \varepsilon.$$

The chain rule is therefore intrinsic to the algebra: any composition of dual-number operations automatically tracks the derivative of the composed function, and no code transformation or symbolic manipulation is required. The same source that evaluates f over $\mathbb{R}$ evaluates both f and $\partial f/\partial x$ when called with dual-number inputs, by virtue of the overloaded operators and elementary functions.

Tangent storage

A dual carries m tangent slots beside its value. They are held in a thread-local bump arena, allocated at the width the call asks for and released in one operation when the solve returns, so m is ncol of the seeding matrix rather than a compile-time constant. A parameter set wider than one call should hold is answered in blocks of directions; since the step sequence of a tangent-blind solve does not depend on m , the blocks share one grid and the assembled Hessian equals a single pass exactly.

An inline slab of compile-time width was the alternative and is no longer used. On a twenty-state chain with forty directions the arena takes 0.86 of its time in the forward mode and 0.74 in forward over forward: a fixed slab is carried whole whether or not a call uses all of it.

One property of the arena form matters for correctness. A dual whose width has never been set has no slots, and reading a tangent from it yields zero rather than faulting. Buffers that are zeroed before use, as every cotangent of a backward sweep is, must therefore have a width declared for them first, or they carry derivatives that are structurally zero.

Sensitivity equations via dual evaluation

Let the ODE right-hand side depend on a parameter vector $p \in \mathbb{R}^{n_p}$,

$$\dot{x} = f(x, p, t), \quad x(t_0) = x_0(p),$$

and define the *sensitivity matrix*

$$S_p = \frac{\partial x}{\partial p} \in \mathbb{R}^{n_x \times n_p}, \quad (S_p)_{ij} = \frac{\partial x_i}{\partial p_j}.$$

To derive the evolution equation for S_p via dual arithmetic, seed each state component with its row of S_p and form the dual-valued state vector

$$\tilde{x}_i = x_i + (S_p)_{ij} \varepsilon_j,$$

where ε_j is the dual unit associated with direction p_j and summation over j is implied. Evaluating f component-wise over the dual algebra and collecting the ε_j -coefficient of the i -th output gives

$$[f(\tilde{x}, p, t)]_{i, \varepsilon_j} = \frac{\partial f_i}{\partial x_k} (S_p)_{kj} + \frac{\partial f_i}{\partial p_j},$$

where repeated k is summed. Since the primal part of the dual evaluation recovers $\dot{x}_i = f_i(x, p, t)$, differentiating $\dot{x}_i = f_i$ with respect to p_j and identifying the ε_j -coefficient with $(\dot{S}_p)_{ij}$ yields

$$\dot{S}_p = \frac{\partial f}{\partial x} S_p + \frac{\partial f}{\partial p}, \quad S_p(t_0) = \frac{\partial x_0}{\partial p}. \quad (11)$$

Hence a single dual-number evaluation of f propagates the full sensitivity matrix S_p in lockstep with the primal state x , at a cost proportional to n_p . No additional Jacobian evaluations are required; the same code path that computes f computes $(\partial f/\partial x) S_p + \partial f/\partial p$ automatically.

Seeding: from S_p to S_θ

In many applications the natural parameters of inference are not the raw entries of p but a higher-level vector $\theta \in \mathbb{R}^M$ to which both the dynamic parameters p_{dyn} and the initial condition $x_0 = p_{\text{init}}(\theta)$ are linked through a user-defined reparametrisation $p = \Phi(\theta)$. The quantity of interest is then

$$S_\theta = \frac{\partial x}{\partial \theta} \in \mathbb{R}^{n_x \times M},$$

related to S_p by the chain rule $S_\theta = S_p \partial p / \partial \theta$. Substituting this relation into (11) yields the evolution equation for S_θ directly,

$$\dot{S}_\theta = \frac{\partial f}{\partial x} S_\theta + \frac{\partial f}{\partial p_{\text{dyn}}} \frac{\partial p}{\partial \theta}, \quad S_\theta(t_0) = \frac{\partial p_{\text{init}}}{\partial \theta}. \quad (12)$$

The matrix $\partial p / \partial \theta$ is the *seed matrix*, supplied by the user (through the `sens1ini` argument of `solveODE()`) and held fixed for the duration of the integration. The choice $\partial p / \partial \theta = I_{n_p}$ recovers the canonical case $S_\theta = S_p$; any other choice carries the sensitivity computation in M specified directions of θ -space at the same cost as integrating M additional states alongside x .

Dual-number aware error norm

When forward sensitivities are active, the integrator carries the sensitivity matrix $S_\theta \in \mathbb{R}^{n_x \times M}$ alongside the primal state. Every accepted step then produces $1 + M$ error vectors of length n_x : the primal estimate $e_n^{(x)} \in \mathbb{R}^{n_x}$ and a sensitivity estimate $e_S^{(k)} \in \mathbb{R}^{n_x}$ for each column $\partial x / \partial \theta_k$ of S_θ . Two natural ways to combine these into a single scalar suggest themselves: a flat WRMS that pools all $n_x(1 + M)$ components into one sum, or a maximum over per-vector WRMS norms. cppDE adopts the second, following the staggered-tolerance convention of CVODES (the `CV_STAGGERED` mode of its sensitivity solver).

The state norm $\|e_n^{(x)}\|_{\text{WRMS}}$ is exactly (9). The sensitivity column norm is its analogue, with S_θ replacing x in both the error vector and the weight,

$$\|e_S^{(k)}\|_{\text{WRMS}} = \sqrt{\frac{1}{n_x} \sum_{i=1}^{n_x} \left(\frac{(e_S)_{ik}}{w_i^{(k)}} \right)^2}, \quad w_i^{(k)} = \text{atol} + \text{rtol} |(S_\theta)_{ik}^{\text{ref}}|, \quad (13)$$

in which $(S_\theta)_{ik}^{\text{ref}}$ is the reference value of the i -th component of column k , taken from the same source (predictor or pre/post maximum) as x_n^{ref} for the state. The scalar that the step-size controller sees is the worst of the state norm and the M column norms,

$$\|e_n\|_{\text{WRMS}} = \max \left\{ \|e_n^{(x)}\|_{\text{WRMS}}, \max_{1 \leq k \leq M} \|e_S^{(k)}\|_{\text{WRMS}} \right\}. \quad (14)$$

The maximum-of-norms convention is preferred to the flat WRMS, which divides by the total component count $n_x(1 + M)$ and therefore averages the squared errors across directions in θ : a large error in one sensitivity column k_0 can be obscured by M small errors in the remaining columns and the state, leading the integrator to accept a step that is too large in the direction k_0 . The maximum norm prevents this dilution and forces each direction in θ to satisfy the requested tolerance independently. Because the controller still operates on a single scalar, the cost relative to the flat norm is a constant overhead in the number of accepted steps, and the control-law parameters $(S, \rho_{\min}, \rho_{\max}, \beta)$ require no adjustment.

The same WRMS structure governs the Newton convergence test in BDF/NDF. A Newton iterate is accepted when the iteration correction satisfies (14), with the maximum again taken over the state and the per-column sensitivity components. Convergence in the primal therefore does not allow a loose solve in any sensitivity direction; a parameter direction that drives the correction norm up triggers either further Newton iterations or a step-size reduction, even when the primal residual is already small.

2.8 Events and restarts

cppDE supports two kinds of events that interrupt smooth integration. A *time-triggered event* fires at a user-specified time $t_e(\theta)$ that may depend on θ through a closed-form expression. A *root-triggered event* fires when a user-supplied event function $g_e(x, t)$ crosses zero along the trajectory; the firing time $t^*(\theta)$ is then defined implicitly by $g_e(x(t^*), t^*) = 0$. In both cases an event applies a discontinuous reset to the state,

$$x^+ = g(x^-, t),$$

typical examples being instantaneous dosing of a kinetic species, a clipping rule that enforces non-negativity, or a switch between dynamical regimes. The remainder of this section first describes the cost of restarting

the integrator after a discontinuity, then sets up the sensitivity-transport problem, and finally derives the correction for each event type.

Restart cost

Single-step methods (Rosenbrock4, Tsit5) carry no information beyond the current state and resume integration in a single full-order step after the reset. Multistep methods (BDF/NDF, Adams) build their order through history and must restart at order 1, rebuilding the Nordsieck array from scratch over several steps. This restart penalty is the dominant cost on problems with frequent events and is the reason Rosenbrock4 is the recommended stiff method when discontinuities are common, even though BDF/NDF is the default elsewhere.

The reset also moves the state onto a different trajectory, so the step size that suited the regime before the event is not a starting point for the one after it. cppDE re-estimates it for every method from the post-event state, with the same initial-step estimator the solve begins with and the remaining interval as its upper bound. For a single-step method that estimate is the whole restart; for a multistep method it comes on top of the order-1 restart above.

Sensitivity transport across a discontinuity

The reset map g acts on the trajectory $x(t, \theta)$ at the firing time $t^*(\theta)$. The naive update of the sensitivity, by the linearisation of g around the pre-event state, $S_\theta^+ = (\partial g / \partial x) S_\theta^-$, is incorrect in the presence of an event because the firing time itself depends on θ . A perturbation δt^* of the firing time induces short integrations of $\dot{x} = f$ on either side of the event, with f taking different values before and after the reset. Denote by f^- the right-hand side just before the reset, evaluated at x^- , and by f^+ the right-hand side just after the reset, evaluated at $g(x^-, t^*)$. The sensitivity at the first grid time following the event then receives two contributions in addition to the linearised reset:

1. *Pre-event drift.* The pre-event trajectory x^- flows along $\dot{x} = f^-$ from the grid time $t_*^{(0)} := \text{scalar}(t^*)$ to the θ -perturbed firing time t^* , over a displacement $\delta t^* = t^* - t_*^{(0)}$ whose scalar part vanishes but whose θ -derivatives do not.
2. *Post-event drift.* The reset state $g(x^-, t^*)$ then flows along $\dot{x} = f^+$ from t^* back to the grid time over the negative of the same displacement.

Both shifts contribute terms of order $f^\pm (\partial t^* / \partial \theta)$ to the sensitivity at the grid time. The correct first-order update reproduces a closed-form expression known as the *saltation matrix*, recovered below as a consequence of the construction.

cppDE realises this correction by replacing the bare reset with a forward-drift-reset-backward-drift sequence, evaluated in dual-number arithmetic. Each drift is integrated by Heun's method, a two-stage second-order Runge-Kutta scheme also known as the improved Euler method, which averages the right-hand side at the start and the end of the interval. Applied to dual-typed input the average is a pure arithmetic operation: the scalar component of x at the grid time is unchanged (since the displacement δt^* has scalar part zero), but the θ -derivatives of the displacement propagate into the state. A single Euler step would suffice for first-order sensitivities, but its truncation error of order δt^{*2} contributes a non-vanishing second-order θ -derivative through the dual algebra and would distort the Hessian. Heun's method has truncation error $O(\delta t^{*3})$ and recovers the correct second-order sensitivities. The reset map g is applied between the two drifts in the same dual-number algebra, so its first- and second-order derivatives propagate automatically.

When AD is disabled (pure-double integration), no transport is required: g is applied to the state as an ordinary function, and the remainder of this section may be skipped.

The construction differs in detail between time-triggered and root-triggered events, because the firing-time perturbation δt enters the dual evaluation differently: directly in the time-triggered case, where $t_e(\theta)$ is supplied by the user and its dual derivatives propagate through the chain rule, and implicitly in the root case, where $t^*(\theta)$ must be reconstructed from the condition $g_e = 0$ via the implicit function theorem.

Time-triggered events

When the firing time is supplied directly as a closed-form expression $t_e(\theta)$, the chain-rule derivatives of t_e with respect to θ are immediately available. The integrator advances on the scalar grid time $t_e^{(0)} := \text{scalar}(t_e)$ and treats the small displacement

$$\delta t = t_e - t_e^{(0)}, \quad \text{scalar}(\delta t) = 0, \quad (15)$$

as a dual number whose value is zero but whose first-order tangents are $\partial t_e / \partial \theta_a$ and whose second-order tangents are $\partial^2 t_e / \partial \theta_a \partial \theta_b$, both inherited automatically from the user-supplied expression for t_e . The state at the grid time $t_e^{(0)}$ is denoted x^- and is itself a dual number that carries the pre-event sensitivities S_θ in its first-order tangent slots.

The transport across the event consists of three steps: a short Heun integration along $\dot{x} = f$ from $t_e^{(0)}$ forward to t_e , the reset map g applied at t_e , and a short Heun integration backward from t_e to $t_e^{(0)}$,

$$\begin{aligned} \text{forward: } & f_1 = f(x^-, t_e^{(0)}), \quad x_E = x^- + f_1 \delta t, \quad f_2 = f(x_E, t_e), \\ & x^* = x^- + \frac{1}{2}(f_1 + f_2) \delta t, \\ \text{reset: } & x_+^* = g(x^*, t_e), \\ \text{backward: } & \tilde{f}_1 = f(x_+^*, t_e), \quad x_B = x_+^* - \tilde{f}_1 \delta t, \quad \tilde{f}_2 = f(x_B, t_e^{(0)}), \\ & x^+ = x_+^* - \frac{1}{2}(\tilde{f}_1 + \tilde{f}_2) \delta t. \end{aligned} \quad (16)$$

The forward leg evaluates the right-hand side at both endpoints of the short interval ($t_e^{(0)}$ and t_e); averaging the two values is the standard Heun update and produces a result accurate to second order in δt . The reset map g is then applied on the event surface at t_e , and the backward leg is the same Heun rule shifted from t_e back to the grid time, this time using the post-event right-hand side.

A single Euler step would suffice for first-order sensitivities: since $\text{scalar}(\delta t) = 0$, the powers δt^k for $k \geq 2$ have zero first-order dual content, and any $O(\delta t^2)$ term drops out of $\partial x^+ / \partial \theta$. For second-order sensitivities, however, the curvature term

$$\frac{1}{2}(Jf + \partial f / \partial t) \delta t^2, \quad J = \partial f / \partial x,$$

acquires a non-vanishing second-order dual component, because the square of δt has a non-zero mixed derivative

$$\frac{\partial^2 (\delta t)^2}{\partial \theta_a \partial \theta_b} = 2 \frac{\partial t_e}{\partial \theta_a} \frac{\partial t_e}{\partial \theta_b},$$

despite both its value and its first derivatives being zero. The Heun midpoint absorbs this curvature term through the second right-hand-side evaluation at the shifted state *and* the shifted time. The use of the dual-typed time argument t_e rather than $t_e^{(0)}$ alone in f_2 and $\tilde{f}_2$ is therefore essential: omitting it would discard the $\frac{1}{2}(\partial f / \partial t) \delta t^2$ contribution for systems with explicit time dependence, even though the time shift itself is invisible at the scalar level ($\text{scalar}(\delta t) = 0$).

Root-triggered events

For a root event the firing time $t^*(\theta)$ is defined implicitly by

$$G(t^*(\theta), \theta) := g_e(x(t^*(\theta), \theta), t^*(\theta)) = 0. \quad (17)$$

Three steps lead from this condition to the corrected state at the next grid time: localisation of t^* on the scalar trajectory, reconstruction of the dual residual δt^* via the implicit function theorem, and the same forward-reset-backward Heun transport as in the time-triggered case.

Localisation The integrator reads g_e at the end of every accepted step and at every requested output time. A crossing is accepted on $[t_n, t_{n+1}]$ when the two readings have opposite signs, and also when $g_e(x_{n+1}, t_{n+1})$ is exactly zero, which is what a root falling on a sampled time looks like. A zero at the lower end is not a new crossing: it was reported when that sample was taken. That convention is also what keeps an event from firing a second time on the surface it has just been applied on, where g_e is zero up to round-off.

The crossing is then refined on the dense output of the accepted step, so no extra right-hand-side evaluations are spent on the localisation itself. The iteration is a false-position step whose interpolation factor is clamped to $[0.1, 0.9]$, which bounds the stagnation of the classical rule while keeping its speed on a nearly linear g_e . It stops when the bracket is narrower than the user-supplied root tolerance, after at most fifty iterations, or as soon as an iterate returns $g_e = 0$ exactly. What is returned as $t_*^{(0)} := \text{scalar}(t^*)$ is the end of the final bracket that still carries the sign g_e had before the crossing. The dual residual $\delta t^* = t^* - t_*^{(0)}$ must now be reconstructed from (17), since a localisation on the scalar trajectory provides no information about how t^* depends on θ .

The dual residual via the implicit function theorem The implicit function theorem yields the derivative of $t^*(\theta)$ directly from the constraint $G(t^*, \theta) = 0$, without explicit inversion, provided G is smooth at the crossing and $\partial G / \partial t \neq 0$ there. Differentiating (17) once with respect to θ and applying the chain rule gives

$$\frac{\partial t^*}{\partial \theta} = -\frac{G_\theta}{G_t} = -\frac{\nabla_x g_e \cdot S_\theta}{\nabla_x g_e \cdot f + \partial g_e / \partial t} \Big|_*, \quad (18)$$

in which the denominator is the total time derivative $\dot{g}_e = \nabla_x g_e \cdot f + \partial g_e / \partial t$ along the trajectory at the crossing. The non-vanishing-denominator condition $\dot{g}_e \neq 0$ is the geometric statement that the trajectory crosses the event surface transversally rather than grazing it; for grazing crossings the firing time itself is not differentiable in θ and a separate fallback applies, see the discussion of simultaneous events below. In dual-number form, (18) is realised as the quotient

$$\delta t^* = -\frac{g_e(x^-, t_*^{(0)})}{\dot{g}_e(x^-, t_*^{(0)})}, \quad (19)$$

where both g_e and $\dot{g}_e$ are evaluated on the dual-typed state x^- at the scalar grid time. The codegen-emitted partials $\nabla_x g_e$ and $\partial g_e / \partial t$ are required only for the assembly of $\dot{g}_e$; the dual chain rule propagates $\nabla_x g_e \cdot S_\theta$ through $g_e(x^-, t_*^{(0)})$ automatically. The scalar residual of the numerator (the bisection localisation tolerance) is set to zero before the division so that the value of δt^* is exactly zero by construction; only its dual derivatives remain.

The first-order dual content of (19) reproduces (18) exactly. At second order the dual quotient misses one term, because g_e is evaluated at the fixed scalar time $t_*^{(0)}$ rather than at the dual-typed time t^* : the quotient rule cannot supply the implicit G_{tt} -contribution that would arise from differentiating g_e in its time slot. Expanding (17) to second order in θ and matching coefficients identifies the missing piece. Writing $\delta t^* = \delta t_{(1)}^* + \delta t_{(2)}^* + \dots$ and using $G(t_*^{(0)}, \theta^{(0)}) = 0$,

$$0 = G_t \delta t_{(2)}^* + \frac{1}{2} G_{tt} (\delta t_{(1)}^*)^2 + G_{t\theta} \delta t_{(1)}^* \delta \theta + \frac{1}{2} G_{\theta\theta} \delta \theta^2,$$

of which the dual quotient supplies the $G_{t\theta}$ - and $G_{\theta\theta}$ -contributions but not the G_{tt} -contribution. The corrected update is therefore

$$\delta t^* \leftarrow \delta t^* - \frac{1}{2} \frac{G_{tt}}{G_t} (\delta t^*)^2, \quad (20)$$

where

$$G_{tt} = \frac{d^2}{dt^2} g_e(x(t), t) \Big|_*$$

is the second total time derivative of g_e along the trajectory. cppDE supplies G_{tt} either through a codegen-emitted analytic callback or, when unavailable, through a forward finite difference of $\dot{g}_e$. Equation (20) touches only the second-order dual components of δt^* ; the value and first-order parts are unchanged.

Forward-reset-backward sandwich With δt^* in hand, the same forward-reset-backward Heun procedure of (16) carries x^- across the event to the grid time $t_*^{(0)}$, with δt replaced by δt^* and the event time t_e replaced by $t_*^{(0)} + \delta t^*$. The five right-hand-side evaluations and the dual-typed time arguments play exactly the same role as in the time-triggered case.

Simultaneous events Two or more root conditions can cross zero within the same accepted step. cppDE localises each crossing independently and groups any that share the same $t_*^{(0)}$ to within the localisation tolerance into a simultaneous-event cluster. For such a cluster the forward Heun shift is performed once to the common event surface, every triggered reset is applied at that surface (the reset maps are assumed to commute), and a single backward Heun shift returns the state to the grid time. The cost is four right-hand-side evaluations independent of the number of simultaneous events. The dual residual δt^* is computed once from the first event in the cluster with non-zero $\dot{g}_e$ at the surface; the remaining events share the same δt^* , since they all sit on the same grid time with the same parameter dependence inherited through x^- .

If every triggered event in the cluster has $\dot{g}_e = 0$ at the surface (a tangential crossing of all event functions), the denominator in (18) vanishes and the implicit function theorem fails to deliver $\partial t^*/\partial \theta$: the firing time is then insensitive to θ at first order. cppDE detects this case by the smallness of $|\dot{g}_e|$ and falls back to applying the reset as a plain map without the saltation transport, which is mathematically correct in this limit because the two saltation contributions $f^\pm (\partial t^*/\partial \theta)$ both vanish.

Conditions switched on by a jump A reset is a jump, not a crossing of the continuous solution. A root condition that a reset makes true is therefore invisible to the sign-change search above, and cppDE reads every condition on both sides of a jump instead. A condition whose sign changes across the jump fires there, one per sweep so that a reset can switch on the next condition, and one that already held before the jump does not fire again.

The switched resets are applied on the surface of the jump, between the forward and the backward shift of the event that caused it, rather than after the transport has returned to the grid time. The composition of the two reset maps is then carried across the discontinuity by a single sandwich, and its sensitivities are those of the same reset written as a time-triggered event at that time, to both orders.

Terminal conditions are excluded from this path. They stop the integration rather than reset a state, which the reset step has no way to report.

Connection to the saltation matrix

The standard treatment of sensitivity transport across a discontinuity yields a closed-form correction, the *saltation matrix*, which generalises the bare reset Jacobian $\partial g/\partial x$ by an additive term accounting for the θ -dependence of the firing time:

$$S_\theta^+ = \frac{\partial g}{\partial x} S_\theta^- + \left(\frac{\partial g}{\partial t} + \frac{\partial g}{\partial x} f^- - f^+ \right) \frac{\partial t^*}{\partial \theta}.$$

The first-order content of the dual-number forward-drift-reset-backward-drift construction reproduces this expression exactly, both for time-triggered events ($\partial t^*/\partial \theta = \partial t_e/\partial \theta$ supplied directly through the user expression for $t_e(\theta)$) and for root events ($\partial t^*/\partial \theta$ computed from (18)). The motivation for adopting the implicit-function and dual-number formulation in cppDE is that it extends naturally to second-order sensitivities, for which the classical saltation-matrix form has no closed-form analogue in the literature. The implementation resides in `cppde_integrate_times.hpp`.

2.9 Reverse-mode sensitivities: the discrete adjoint

Forward sensitivities cost one additional state vector per direction in θ , so a gradient of a scalar objective over M parameters costs M times the value alone. When the quantity of interest is a single scalar J , a likelihood or a residual sum of squares, the transpose of that computation delivers the whole gradient at a cost independent of M . cppDE emits this second differentiation mode under `cppODE(..., derivMode = "reverse")`.

The quantity computed is a directional contraction rather than the sensitivity matrix. Given a cotangent $W \in \mathbb{R}^{n_t \times n_x}$ attached to the observed trajectory, the reverse mode returns

$$\bar{\theta}_j = \sum_{o=1}^{n_t} \sum_{i=1}^{n_x} W_{oi} \frac{\partial x_i(t_o)}{\partial \theta_j}, \quad j = 1, \dots, n_x + n_p, \quad (21)$$

in one sweep. The row index o runs over the observation times and the column index i over the state components. Write $w_o \in \mathbb{R}^{n_x}$ for the vector formed from row o , the cotangent attached to the trajectory at t_o . For a least-squares objective W is the residual weighted by the error model, and $\bar{\theta}$ is $\nabla_{\theta} J$ exactly. Several cotangents may be supplied as columns of a three-dimensional array, one sweep each.

The adjoint equation

The contraction in (21) can be evaluated without forming S_p at all. The construction is classical and is given here because both backends implement it, each with a different discretisation. Let S_p be the forward sensitivity of (11). Define $\lambda(t) \in \mathbb{R}^{n_x}$ backwards from the end of the observed interval by

$$\dot{\lambda} = -J(x, p, t)^{\top} \lambda, \quad \lambda(T^+) = 0, \quad \lambda(t_o^-) = \lambda(t_o^+) + w_o, \quad (22)$$

so that λ evolves continuously between observation times and jumps by the cotangent row at each one. Along the solution

$$\frac{d}{dt} (\lambda^{\top} S_p) = -\lambda^{\top} J S_p + \lambda^{\top} \left(J S_p + \frac{\partial f}{\partial p} \right) = \lambda^{\top} \frac{\partial f}{\partial p}, \quad (23)$$

where the two Jacobian terms cancel. That cancellation is what allows λ to be defined without reference to θ , and it is the reason the construction is independent of the number of parameters. Integrating (23) over $[t_0, T]$ and collecting the jumps, whose contributions are exactly the terms of (21), gives

$$\bar{\theta} = \left(\frac{\partial x_0}{\partial \theta} \right)^{\top} \lambda(t_0) + \left(\frac{\partial p}{\partial \theta} \right)^{\top} \int_{t_0}^T \left(\frac{\partial f}{\partial p} \right)^{\top} \lambda(t) dt. \quad (24)$$

Taking θ to be the initial state and the parameters themselves makes both outer factors identities and splits the answer in two,

$$\bar{x}_0 = \lambda(t_0), \quad \bar{p} = \int_{t_0}^T \left(\frac{\partial f}{\partial p} \right)^{\top} \lambda(t) dt, \quad (25)$$

which is the pair a reverse solve reports. A reparametrisation $p = \Phi(\theta)$ is applied afterwards through the outer factors of (24), which is the chain rule of (12) applied on the left instead of on the right.

The gradient is therefore obtained from one backward solve of a linear system of the dimension of the state, together with a quadrature of the dimension of p . Neither (22) nor (25) involves M , since $\partial f / \partial p$ enters only through its contraction with λ and is never formed as a matrix. The two parts of (25) are returned stacked in the order in which a forward seed matrix is indexed, states first and parameters second.

Discrete and continuous adjoints

Equation (22) describes the exact trajectory. Two discretisations of it are in common use, and `cppDE` provides both.

The *continuous* adjoint differentiates first and discretises afterwards: (22) is handed to an integrator of its own and solved backwards over forward states recovered by interpolation [6]. This is the route `CVODES` takes, and Section 3.3 describes the form in which `cnode()` exposes it.

The *discrete* adjoint reverses that order. It transposes the map the forward run applied, one accepted step at a time, and is what `cppODE()` implements. Both approximate the same gradient, and their results converge as the tolerance is lowered, but they differ in three respects that matter in practice.

Consistency of value and gradient. The discrete adjoint differentiates the trajectory that was actually computed, so it returns the derivative of the value the forward run reported, up to rounding. The continuous

adjoint approximates the derivative of the exact solution, which no run computes, so its gradient and the accompanying value are consistent only to $O(\text{tol})$. Tightening the backward tolerance alone does not remove that difference.

Discontinuities. A jump in the state implies a jump in λ , given by the transposed saltation matrix of the event. A continuous adjoint has to restart its backward integration with that matrix, which requires an interface through which the solver can be informed of the discontinuity. In the discrete adjoint the jump is the transpose of arithmetic the forward run performed, and is handled by the same replay as every other step.

Linear algebra per step. The adjoint of an implicit step is a single transposed solve against the iteration matrix of that step, evaluated at the state stored in its checkpoint, without a Newton iteration and without error control of its own. The continuous adjoint integrates a stiff linear ODE instead, with a step-size controller and a step count of its own that need not resemble those of the forward run.

The step map and its adjoint

Reverse accumulation cannot be obtained from an augmented arithmetic in the way forward sensitivities are. A dual number carries its derivative alongside its value, so the derivative of an operation is available as soon as the operation has executed. An adjoint is a sum over all later uses of a quantity, none of which has occurred at that point in the computation.

Two implementations are possible. The first records the operations of a step and traverses the record backwards; cppDE used this until 2026-09-11. The second uses the fact that the map to be differentiated is not arbitrary. It is one accepted step of a known method on a grid that is held fixed, so its adjoint can be written in closed form. cppDE uses the second implementation.

An accepted step of the multistep family decomposes into three parts:

$$z^{\text{pred}} = A z^{\text{in}}, \quad (26)$$

$$r(y, z^{\text{pred}}, \theta) = 0, \quad (27)$$

$$z^{\text{out}} = B z^{\text{pred}} + c a, \quad a = y - z_0^{\text{pred}}, \quad (28)$$

namely the Nordsieck rescale and Pascal shift, the corrector, and the tail that the controller applies between the acceptance of this step and that of the next. The operators A , B and c act on the slot index only; their coefficients follow from the carry and are identical for every state component. They are therefore not transposed by hand but read off the stepper: a probe whose states are the Nordsieck slots is driven through the routines the forward step runs, carrying the identity in the $[\text{slot} \times \text{component}]$ array, and one pass returns every column. Since no coefficient is restated, the adjoint cannot drift from the step it belongs to.

The corrector is differentiated by the implicit function theorem, which is what makes the adjoint of an implicit step inexpensive. The derivative $\partial r / \partial y$ is the iteration matrix $W = I - \gamma J$ that the forward step factorises, so one transposed solve against W replaces the entire Newton iteration. For the Rosenbrock family the corresponding object is a transposed stage recursion, six solves against one factorisation; for an explicit method it is the backward recursion over the tableau.

The contractions the adjoint requires are supplied by the model and emitted by the code generator beside the right-hand side: $J^\top \lambda$, $(\partial f / \partial p)^\top \lambda$ and, for a Rosenbrock method, the same two applied to Jv , which carries a second derivative because its stages solve against a matrix built from the Jacobian.

Checkpoint per step

The backward pass requires the states the forward run passed through. An adaptive run takes hundreds to thousands of steps, so these are stored not as a tape of operations but as one checkpoint per accepted step, holding the step-start state, the carry, the time, the step size the controller chose, and a record of the operations the controller applied to the carry afterwards. The memory required is linear in the number of steps and contains no derivative information.

Observations lie at interpolated times rather than at step ends, and are evaluated through the continuous extension of the step that contains them. The adjoint of that evaluation is obtained from the same probe: after the tail has been applied the probe holds a valid interpolant, and evaluating it at the observation time yields the row $\partial x_{\text{interp}}/\partial(z^{\text{pred}}, a)$ directly.

Checkpoint reuse

The sweep reads nothing but the checkpoints, so the forward pass and the backward pass can be issued as separate calls. A reverse solve with `keepStore = TRUE` returns its checkpoints as `$store`, and a later solve of the same model at the same times and parameters accepts them through `store`; that solve performs no integration and proceeds directly to the sweep. A store created at a different parameter value is rejected rather than reused. The seed may be omitted on the call that creates a store, which then performs a value-only run and retains its checkpoints.

A gradient requires two solves at the same parameter value, one to provide the values from which the seed is formed and one to perform the sweep, and the store removes the second integration. Several cotangents are handled in the same way: the columns of a three-dimensional seed sweep the one store in turn, so n_{seed} cotangents cost n_{seed} sweeps and a single integration.

Table 4 compares the routes on one objective chain, `Bachmann_MSB2011` from the collection of Section 5, with $n_x = 25$ over 36 experimental conditions, including the observation map and the residual sum of squares, and differentiated in $M = 113$ parameters.

Route	ms	in value runs
Value only	34.8	1.00
Forward sensitivities	974.6	28.00
Discrete adjoint	106.6	3.06
CVODES adjoint sensitivities	201.1	5.77

Table 4: One gradient of one objective chain, $n_x = 25$ over 36 conditions, $M = 113$, at fixed tolerances on one core. The value run is the same chain without derivatives. The forward column is the cost the adjoint is meant to remove, and grows with M where the other two do not.

Per step, the backward pass costs about 1.7 times the forward integration, measured over the 721 accepted steps of the longest condition. Roughly half of that is linear algebra: the sweep re-forms and refactorises the iteration matrix at each checkpoint, the factorisation of the forward run having been discarded long before the sweep reaches that step. The slot operators are cached rather than rebuilt, keyed on the carry and the step size, so the probe described above is invoked only where the controller changed the control state, which occurs at roughly one step in five.

The adjoint of an event

An intervention is a map of its own between two steps. The engine transports the discontinuity with a Heun sandwich: a shift to the event surface, the resets, and a shift back to the grid time. Each shift is taken over the residual of the event time, whose own value is subtracted before it is used, so the residual is numerically zero and only its derivative survives. In value the jump is therefore a plain jump, and the sandwich reduces to the classical saltation relation,

$$dx = R'(dx_{\text{before}} + f_{\text{before}} s) - f_{\text{after}} s, \quad (29)$$

with R the resets, R' their Jacobian and s the differential of the event time. For a root event $g(x, t, p) = 0$ the implicit function theorem gives $s = -(\nabla g \cdot dx + \partial g/\partial p \cdot dp)/\dot{g}$, and for a fixed event at $t = \tau(p)$ it is $s = (\partial \tau/\partial p) \cdot dp$.

Since (29) is linear in dx_{before} and dp , its transpose can be read off directly. Contracting it with a cotangent

λ_{after} and collecting the terms that multiply s leaves one scalar,

$$\kappa = \lambda_{\text{after}}^\top (R' f_{\text{before}} - f_{\text{after}}), \quad \lambda_{\text{before}} = R'^\top \lambda_{\text{after}} + \kappa \left(\frac{\partial s}{\partial x} \right)^\top, \quad (30)$$

and the parameter cotangent receives $\kappa (\partial s / \partial p)^\top$ in addition. For a root event both derivatives of s carry the factor $-\kappa / \dot{g}$, so the state picks up $-(\kappa / \dot{g}) \nabla g$ and the parameters $-(\kappa / \dot{g}) \partial g / \partial p$. For a fixed event $\partial s / \partial x = 0$ and the parameters pick up $\kappa \partial \tau / \partial p$.

The transpose $R'^\top$ is applied reset by reset and in reverse order. A reset that replaces a state removes the incoming cotangent of that component, one that adds to it leaves the cotangent unchanged, and one that multiplies it by h scales it by h ; where the reset value h depends on the state or on the parameters, its derivatives enter the same transpose. The adjoint of a jump therefore requires the right-hand side at the two ends, ∇g and $\dot{g}$, and the model's derivatives of the reset and event-time expressions. It requires no Jacobian of f , because every term the two shifts would contribute carries the residual as a factor and that factor is zero.

After a jump the multistepper restarts. `initialize` builds the Nordsieck history from a single state as $z_0 = x$ and $z_1 = h f(x, t)$, with the remaining slots zero, so the cotangent of the history collapses onto that one state. Writing w_j for the cotangent of slot j , the restart contributes

$$\bar{x} = w_0 + h J^\top w_1, \quad \bar{p} = h \left(\frac{\partial f}{\partial p} \right)^\top w_1, \quad (31)$$

and slots above the first contribute nothing. The start of the trajectory applies the same map, and both are served by one routine.

Second order of an event's adjoint

Write s for the differential of the event time, so that $s = 0$ in value and $ds \neq 0$. Then

$$s^2 = 0, \quad d(s^2) = 2s ds = 0, \quad d^2(s^2) = 2(ds)^2 \neq 0, \quad (32)$$

so terms of order s^2 drop from the gradient and not from the Hessian. Three of them enter a jump. The same three enter a fixed event, with the event time's own residual $\tau = \varepsilon(\theta) - \lfloor \varepsilon \rfloor$ in place of s : a fixed event is the same five-stage sandwich, and its transpose is (??) with τ substituted throughout. Carrying only the classical saltation there is exact in value and wrong in the tangents as soon as the jump height moves with a parameter, because $\bar{x}_a = (I - \tau J_a)^\top \bar{x}_o$ then has a second factor that is not zero in value.

The shift. Each half of the sandwich is a Heun step,

$$x_* = x_{\text{before}} + f_{\text{before}} s + \frac{1}{2} a_{\text{before}} s^2, \quad a = \frac{\partial f}{\partial x} f + \frac{\partial f}{\partial t}, \quad (33)$$

and $a s$ is the difference of the two right-hand sides the step already evaluates, so no Jacobian is needed.

The reconstructed state. The store keeps x_{before} alone. Rebuilt from the bare reset, the post-jump state satisfies $\partial x_{\text{after}} / \partial t^* = 0$, whereas

$$\frac{\partial x_{\text{after}}}{\partial t^*} = f_{\text{before}} - f_{\text{after}}. \quad (34)$$

The reconstruction therefore runs the sandwich, not the reset.

The quadrature. The parameter cotangent

$$\bar{p}_k = \int_0^T G_k dt = \int_0^{t^*} + \int_{t^*}^T, \quad G_k = \left(\frac{\partial f}{\partial p_k} \right)^\top \lambda, \quad (35)$$

is split at t^* , and G_k is discontinuous there because the state is. Differentiating a split integral whose split point moves leaves the jump of the integrand at the split,

$$\frac{\partial \bar{p}_k}{\partial v} \supset \left(G_k(t^{*-}) - G_k(t^{*+}) \right) \frac{\partial t^*}{\partial v}. \quad (36)$$

Only parameters whose $\partial f/\partial p_k$ reads a state the reset changed contribute.

Implemented as accumulations scaled by the residual, all three are zero in value and carry $\partial t^*/\partial v$ in their tangent, so the gradient is unchanged.

For a fixed event at $t = \tau(p)$ the residual is $\tau - \lfloor \tau \rfloor$. For a root event it is

$$s = -\frac{g - \lfloor g \rfloor}{\dot{g}}, \quad \dot{g} = \frac{\partial g}{\partial t} + \nabla g \cdot f, \quad (37)$$

read at the stored state. Localisation is tangent-blind, so that state carries $\partial x/\partial \theta$ at a fixed time and ∇g contracts it to the event time's.

Transposing the root sandwich

The jump is the composition

$$x_e = x_b + f_b s, \quad x_* = x_b + \frac{1}{2}(f_b + f_e)s, \quad x_a = R(x_*), \quad x_k = x_a - f_a s, \quad x_o = x_a - \frac{1}{2}(f_a + f_k)s, \quad (38)$$

with $f_e = f(x_e, t^*)$, $f_a = f(x_a, t^*)$ and $f_k = f(x_k, t_0)$. Equation (38) is a straight-line program, and its adjoint is obtained by reversing it assignment by assignment rather than by collecting a closed form. A closed form invites the shift to be pruned by order in s , retaining a product only where one factor does not vanish in value. That rule is valid for the gradient and invalid for the Hessian: by (32), s^2 has vanishing value and vanishing first derivative, while $d^2(s^2) = 2(ds)^2 \neq 0$.

Reversed assignment by assignment, the shift receives a contribution from every assignment in which it occurs. Two of these arise because f_e and f_a are evaluated at t^* , which depends on θ :

$$\bar{s} \supset \left(\frac{\partial f}{\partial t} \right)_{(x_a, t^*)} \cdot \bar{f}_a + \left(\frac{\partial f}{\partial t} \right)_{(x_e, t^*)} \cdot \bar{f}_e. \quad (39)$$

Both cotangents carry a factor s , so (39) is absent from the gradient and contributes at full strength to the Hessian. It vanishes identically for an autonomous right-hand side, and a pruned form omits it precisely when $\partial f/\partial t \neq 0$.

The forward path corrects the shift by $-\frac{1}{2}(\ddot{g}/\dot{g})s^2$, whose chain factor $1 + 2c_2s$, $c_2 = -\ddot{g}/(2\dot{g})$, is one in value and a second-order tangent; the transpose carries it on $\partial s/\partial x$.

Three checks bound the implementation. With $R = \text{id}$ the composition (38) is the identity and the transpose cancels term by term. `dev/jump-adjoint-check.cpp` compares the transpose with a nested-dual evaluation of the same sandwich, without a solver, a store or an event engine. `dev/cxx/run.sh --reverse-events2` compares a whole trajectory with forward over forward on one recorded step sequence. Only the third is sensitive to (39): the second is checked against the same forward routine the transpose reverses, so a term absent from both sides cancels.

The shift has a derivative of its own. Since $\dot{g}$ reads the state,

$$\frac{\partial s}{\partial x} = -\frac{\nabla g}{\dot{g}} - \frac{s}{\dot{g}} \nabla \dot{g}, \quad \nabla \dot{g} = \nabla \frac{\partial g}{\partial t} + (\nabla^2 g) f + J^\top \nabla g, \quad (40)$$

and the second term carries g , which is zero in value: absent from the gradient and full strength in the Hessian. For a g that is linear in the state and blind to the time, $\nabla \dot{g}$ is $J^\top \nabla g$ alone, which is what makes that case a poor test of it. The generator therefore differentiates $\dot{g} = \partial g/\partial t + \nabla g \cdot f$ as one expression rather than assembling three terms, which gives $\nabla \partial g/\partial t$ and $(\nabla^2 g)f$ with it, and on θ the matching $\partial^2 g/\partial \theta \partial t$, $(\partial \nabla g/\partial \theta) \cdot f$ and $(\partial f/\partial \theta)^\top \nabla g$. A forcing inside g cannot be differentiated in t by the algebra system, so its rate enters as a slot that prints as the interpolant's own derivative.

A reset that depends explicitly on time. The value h is applied on the surface, at t^* , so

$$\frac{\partial x_a}{\partial t^*} \supset c \frac{\partial h}{\partial t}, \quad c = \begin{cases} 1 & \text{add} \\ 1 & \text{replace} \\ x_{*k} & \text{multiply,} \end{cases} \quad (41)$$

which contributes to κ , the scalar multiplying the shift. This is a first-order term: omitting it makes the gradient wrong, not only the Hessian. For the same reason the transposed reset receives the event time as the differentiable type rather than as a plain number, since otherwise $\partial h/\partial\theta$ cannot carry the dependence of t^* on θ .

Such a value inherits the root tolerance rather than the integration tolerance. It evaluates t^* directly, whereas a value depending only on the state at t^* sees the localisation error damped by the smoothness of the trajectory.

Against the sandwich itself, evaluated over a nested dual with no solver around it, the transpose is exact to $9 \cdot 10^{-16}$, and over a whole trajectory the second order agrees with the forward mode to the tolerance of the solve.

Reaching that required the contraction buffers to be bound by the caller. A generated $J^\top\lambda$ opens an arena scope of its own and writes its result inside it, so an output the sweep has not armed retains a tangent pointer into storage that the scope reclaims on exit. The buffer for $J^\top\nabla g$ is filled twice, and κ and κs are formed between the two fills, within the reclaimed window; the second fill overwrote their tangents. Every read of the buffer is multiplied by s , which vanishes in value, so the lost tangents were immaterial and the corruption of κs was not: it produced a rank-one error on $(J^\top\nabla g, (\partial f/\partial\theta)^\top\nabla g)$ alongside a bit-identical gradient. Fixed events do not call the contraction and were unaffected.

The failure mode is worth recording, because an incorrect Hessian of this kind is easily misread. The gradient was exact, every shift factor vanishes in value, and the residual was rank one on a vector the derivation does contain, so the symptom is indistinguishable from a missing analytical term. It was separated from one by a check that alters no algebra: the same transpose evaluated once with buffers owned by the sweep and once with buffers bound as the generated code binds them.

Verifying a derivative across a moving event

When t^* depends on a parameter, the reported trajectory carries a row whose time moves with it. A functional $w^\top x$ seeded on that row is not the same functional at two parameter values and has no derivative to compare against. Restricting the seed to the user's own output times restores a well-posed functional, after which the forward, forward-over-forward and difference-quotient answers agree to the tolerance of the solve.

Internal numerical differentiation

The step sizes the controller selected are constants of the map that is differentiated. This is Bock's internal numerical differentiation: the nominal run adapts freely, and the derivative is taken of the scheme that run applied. Differentiating the control law instead introduces derivatives of the step sizes into the chain rule, which makes the discrete adjoint inconsistent with the adjoint ODE. That term was measured at $O(\text{tol})$ and is deliberately absent.

The choice is often defended by the claim that h is piecewise constant in θ and its derivative therefore zero almost everywhere. That is not what an adaptive controller produces, and the argument below decides the question without it.

Let $y_h(\theta)$ denote the result the program returns on the grid the controller selects at θ . Then

$$\frac{dy_h}{d\theta} = \frac{\partial y_h}{\partial\theta} + \sum_k \frac{\partial y_h}{\partial h_k} \frac{dh_k}{d\theta}, \quad (42)$$

where the first term is what internal numerical differentiation computes and the sum is the grid effect, which differentiating the control law would add.

The function $h(\theta)$ is piecewise smooth with jumps. Within an acceptance region it varies smoothly through the error estimate; at the boundary of one a step that was accepted is rejected and the entire later grid shifts. Halving the spacing of a sweep in θ doubles the measured magnitude of the grid effect, which identifies it as a jump height divided by $\Delta\theta$ rather than as a derivative. Automatic differentiation reproduces only the smooth part. Differentiating the control law therefore returns the exact gradient plus the smooth part of the

grid effect, with the jump part, which is the larger of the two, absent. The result is neither the derivative of the exact solution, to which the addition does not belong, nor the derivative of the reported value, of which the dominant contribution is missing. The smooth part is a correct derivative, but not of either quantity of interest. Recovering the grid effect requires a smoothed controller, a fixed grid or a randomised estimator, none of which differentiation of the existing chain provides.

By the same criterion it is the forward mode that departs from internal numerical differentiation, not the reverse one. The value pass of the reverse mode instantiates the stepper on plain numbers, so the controller reads the state error alone and the derivative is taken on the grid of the nominal run. The dual forward mode maximises its error norm over the sensitivity columns as in (14) and therefore selects its grid from value and tangents together, integrating on a finer grid to which no reported value belongs. The reverse gradient is thus the derivative of the trajectory a pure value run returns, and the forward gradient the derivative of a trajectory that is not reported; the $O(\text{tol})$ difference between them is attributable mainly to the forward mode.

The sensitivities remain inside the error norm in either case. They are integrated alongside the state, their errors accumulate, and omitting them from the control is the situation `CVodeSetSensErrCon` exists to warn against. Value and gradient on slightly different grids is the smaller of the two costs. `solveODE(..., sensErrCon = FALSE)` selects the alternative.

Goal-oriented step-size control

A grid selected from the state error alone need not be fine enough for the gradient. Since λ is solved backwards, a step at which λ is large contributes strongly to the error in J even where its own state error is small; the gradient inherits the order of the state error but not its constant. In the forward mode this is settled within the run, by maximising the error norm over the sensitivity columns as in (14). In the reverse mode it cannot be, because λ at step k depends on every step after it and is not available while step k is being taken.

It is available across runs, and the `errWeights` argument of `solveODE()` accepts it: the λ left by an earlier sweep, sampled on that sweep's grid, is read back by linear interpolation and entered into the error norm as

$$\text{err} = \max \left(\text{err}_{\text{state}}, \frac{|\lambda(t_{n+1})^\top e_n|}{\text{gradtol}} \right), \quad (43)$$

where e_n is the embedded error estimate the controller forms in any case. The numerator estimates the contribution of this step to the error in the objective, and `gradtol` is the tolerance to which that contribution is held. Without a tolerance of its own the two terms of (43) would not be comparable.

The use of a maximum has two consequences. The grid can only become finer than `abstol` and `reltol` require, never coarser, so the accuracy of the trajectory is unaffected and the weighting is safe to leave enabled. And since a weight can only reduce a step size, an inaccurate λ costs at most computing time: if it is too large the steps are unnecessarily small, and if it is too small the state term governs as before. An estimate of λ taken at a neighbouring parameter value is therefore admissible, which is the case of practical interest, since an optimiser can pass the sweep of iteration $i - 1$ to iteration i and pay for one pass rather than two.

The weight is an input to the controller and is held in plain `double`. The controller is not differentiated, and a weight carried into the derivative would reintroduce the step-size term discussed above. The interpolation is linear rather than PCHIP for the same reason and for two further ones: λ jumps where the objective seeds it and where an event resets the state, so `breaks` names the sample indices an interpolant must not span, and a query inside such an interval takes the nearer endpoint; and shape preservation is of no value for a weight whose overshoots cost only time. A `floor` raises weights that fall below a given fraction of the largest, which is a precaution rather than a necessity, since λ measures a linearised influence and a state it weights near zero may still move nonlinearly.

The indicator can be inspected directly. With `adjointGrid = TRUE` a reverse solve also returns the grid the sweep ran on, λ at the start of each step, and $\eta_n = \lambda(t_{n+1})^\top e_n$ for each step, which is the numerator of (43) and estimates that step's contribution to the error in J . The estimate is of first order, since it evaluates λ at

the computed rather than the exact state, and what the adaptation requires of it is the ranking of the steps rather than the magnitude.

Relation to the forward mode

Forward and reverse derivatives of the same model do not agree to machine precision. A sensitivity run maximises its error norm over the state error and every tangent column, so at equal tolerances its controller selects a finer grid than a value-only run. The two modes therefore differentiate two different discretisations, each of them exactly, and the difference between the results falls with the tolerance. On the chain of Table 4 the angle between the two gradients is 3.6×10^{-7} radians.

One consequence favours the reverse mode independently of its cost. Its gradient is the derivative of the trajectory that a value-only solve produces, so the value and the gradient returned to the caller are consistent with each other. This is of practical importance in a trust-region method, which compares a predicted reduction with an achieved one.

Of the four methods, the Rosenbrock adjoint is the most accurate, for a structural reason: its stages are direct solves, so its adjoint is the exact transpose of the discrete step, whereas the adjoint of a corrector is the implicit function theorem applied to an iteration that was stopped at a finite tolerance.

The reverse mode supports all four time-stepping methods, time- and root-triggered events, forcings, the batch entry point and the sparse linear solver. The implementation resides in `cppde_adjoint_step.hpp`, `cppde_reverse_step.hpp` and `cppde_reverse_trajectory.hpp`.

Second derivatives: forward over reverse

For a fixed cotangent, (21) defines the gradient map

$$\bar{\theta} : \mathbb{R}^M \rightarrow \mathbb{R}^M, \quad \theta \mapsto \nabla_{\theta} J(\theta), \quad (44)$$

whose Jacobian is the Hessian $\nabla_{\theta}^2 J$. Nothing in the sweep that evaluates (44) depends on the scalar type, so instantiating it on `dual(double, N)` rather than on `double` differentiates the sweep itself. Seeded with directions $V \in \mathbb{R}^{M \times N}$ it returns the pair $(\nabla_{\theta} J, \nabla_{\theta}^2 J V)$, both computed exactly rather than by a difference quotient; for $V = I_M$ the second member is the Hessian. This is selected by `derivMode = "forward-reverse"` and reported as `adjoint2`, of shape $[n_x + n_p, N, n_{\text{seed}}]$, beside the `adjoint` and `sens1` of the same run.

What the tangents satisfy is the second-order adjoint system. Write $v = (v_0, v_p)$ for one direction in θ , let σ be the directional sensitivity and μ the tangent of λ along v . Then σ solves (11) contracted with v ,

$$\dot{\sigma} = \frac{\partial f}{\partial x} \sigma + \frac{\partial f}{\partial p} v_p, \quad \sigma(t_0) = v_0, \quad (45)$$

and differentiating (22) along v gives μ backwards from $\mu(T^+) = 0$,

$$\dot{\mu}_i = - \sum_k \frac{\partial f_k}{\partial x_i} \mu_k - \sum_k \lambda_k \left(\sum_j \frac{\partial^2 f_k}{\partial x_i \partial x_j} \sigma_j + \sum_m \frac{\partial^2 f_k}{\partial x_i \partial p_m} v_m \right), \quad (46)$$

with the same jumps as λ , which vanish here because the cotangent rows do not depend on θ . Differentiating (25) along v then gives the two halves of $\nabla_{\theta}^2 J v$, the state half as $\mu(t_0)$ and the parameter half as

$$\int_{t_0}^T \left(\sum_k \frac{\partial f_k}{\partial p_m} \mu_k + \sum_k \lambda_k \left(\sum_j \frac{\partial^2 f_k}{\partial p_m \partial x_j} \sigma_j + \sum_l \frac{\partial^2 f_k}{\partial p_m \partial p_l} v_l \right) \right) dt. \quad (47)$$

cppDE does not integrate (45) to (47). It obtains their discrete counterpart from the dual instantiation: the tangent carried alongside λ is μ on the grid the forward run chose, and the tangent of the parameter accumulator is (47). The second derivatives of f that appear in (46) and (47) are supplied by differentiating

the first-order contractions $J^\top \lambda$ and $(\partial f / \partial p)^\top \lambda$ in dual arithmetic, so the code generator emits no second-derivative expression of its own.

The tangents have to be present in the forward pass as well, since the sweep reads its checkpoints, which makes that pass a sensitivity run in the same N directions. The cost is one sensitivity integration and one sweep, both of width N , and therefore linear in N . Forward over forward, selected by `derivMode = "forward-forward"` or by the older `deriv2 = TRUE`, propagates the full second-order sensitivity tensor through the integration at a cost quadratic in M . A difference quotient of the gradient would cost N further gradient evaluations and lose half of the available digits.

The quadratic route is not superseded, because it answers a different question. It returns `sens2`, the second derivatives of every state with respect to every pair of parameters, from which the Hessian of any functional can be contracted afterwards. Forward over reverse returns the Hessian of one functional. Where that functional is a likelihood the mixed route is the cheaper by a factor of M ; where the states themselves are of interest, only `sens2` answers.

Two facilities are required that the first order does not need. The iteration matrix now carries tangents, so the transposed solve applies the implicit function theorem with the matrix transposed alongside; `cppde_ad_lu.hpp` and `cppde_sparse_ad_lu.hpp` provide this in the dense and the sparse case. And a generated model body opens an arena scope of its own, in which a tangent allocated for one of the caller's numbers would be destroyed, so the sweep binds the tangent storage of every buffer before passing it to the model, as the stepper does for its own buffers through the tangent slab.

Verification is against the forward mode. The two second-order routes agree to 10^{-6} on `bdf`, `rb4` and `tsit5`, and all four methods return the same Hessian backwards to 10^{-5} . The second comparison is the stronger one, since the four methods adapt grids of their own, so agreement implies that each sweep differentiated the discretisation its own method applied.

2.10 Method selection

The recommended choice of method for the most common scenarios is summarised in Table 5.

Situation	Recommended method
Default; stiffness unknown	"bdf"
Stiff, frequent events	"rb4"
Stiff, smooth, large n_x with sparse Jacobian	"bdf" with <code>sparse = TRUE</code>
Non-stiff, expensive right-hand side	"tsit5"
Non-stiff, cheap right-hand side, long horizon, high accuracy	"adams"

Table 5: Recommended `method` argument for common integration scenarios. Stiffness is the dominant factor; in its absence, right-hand-side cost and required accuracy decide.

The cost of an inappropriate choice is asymmetric. An explicit method applied to a stiff problem will fail or be driven to step sizes that make the integration uneconomic, while an implicit method applied to a non-stiff problem incurs only a constant overhead per step. In the absence of prior knowledge of the spectrum of J , BDF/NDF is the safer default.

3 `cvode()`

The `cvode()` entry point compiles a model against the system-installed SUNDIALS CVODE/CVODES library and exposes the result through the same `solveODE()` interface as `cppODE()`. The two backends therefore differ only in the numerical core; the codegen pipeline, the `events`, `forcings`, `rootfunc`, and `fixed` arguments, the shape of the returned object, and the column ordering of `parms` and `sens1ini` all behave identically. SUNDIALS (≥ 6.0) must be installed on the build host; KLU (used for sparse Jacobians) is detected the same way and is required when `sparse = TRUE`.

3.1 What is shared

The symbolic Jacobian $J = \partial f / \partial x$ is derived through SymPy at code-generation time and emitted as analytic C++ exactly as in `cppODE()`. The KLU sparse linear solver path activates automatically when the Jacobian is sufficiently sparse and the SuiteSparse libraries were detected at install time; the dense fallback uses LAPACK. Forcing functions are PCHIP-interpolated, events apply a discontinuous reset to the state, and root events fire when a user-supplied scalar function crosses zero.

3.2 What differs

Sensitivity right-hand side. CVODES integrates the analytic forward-sensitivity equation

$$\dot{S}_k = J(x, p) S_k + \frac{\partial f}{\partial p_k}, \quad k = 1, \dots, n_p, \quad (48)$$

directly. The in-tree backend recovers the same equation by evaluating f on dual-typed state, cf. (11). The two formulations are mathematically equivalent. Numerically, the analytic right-hand side avoids the per-evaluation arena allocation required by the dual-number path and is preferable on large problems with a sparse $\partial f / \partial p$; conversely, the dual-number path benefits from SymPy’s common-subexpression elimination on small problems for which integrator control flow dominates the cost.

Available methods. CVODE provides only the multistep families. Calling `cvode(method = "bdf")` or `cvode(method = "adams")` selects between them; the single-step methods `rb4` and `tsit5` are not available through this backend. The cost-per-step accounting and the stability properties documented for the multistep methods in Section 2 therefore carry over directly.

Second-order sensitivities. CVODES exposes only first-order forward sensitivities; the nested-dual construction has no counterpart in the SUNDIALS API. `cvode()` therefore does not accept `deriv2 = TRUE`.

Saltation transport across events. Events on `cvode()` models use the same forward-drift-reset-backward-drift sequence as `cppODE()`, with the first-order saltation correction of equations (16) and (18). The transport is performed in the generated glue layer rather than inside CVODES, since CVODES exposes no entry point at which the sensitivity solver can be informed of the discontinuity.

Reparametrisation seeding. The chain-rule construction $S_\theta = S_p \partial p / \partial \theta$ of equation (12) is identical to the `cppODE()` case. `solveODE()` accepts the same `sens1ini` matrix shapes (legacy, full, and partial) and the same row-name conventions; only second-order S_θ is unavailable.

3.3 Adjoint sensitivity analysis

`cvode(..., derivMode = "reverse")` compiles the CVODES adjoint sensitivity analysis (ASA) in place of the forward sensitivities and returns the same `$adjoint` as a native reverse solve, namely the contraction (21) of a seed, stacked states first and parameters second. The remainder of the call is unchanged, including the batch entry point.

The backward problem. Where `cppODE()` transposes the steps the forward run took, CVODES integrates (22) as an ODE in its own right [6]. The generated glue emits the two transposed contractions from the symbolic Jacobian the forward path already uses,

$$\dot{\lambda} = - \left(\frac{\partial f}{\partial x} \right)^\top \lambda, \quad \dot{q} = - \left(\frac{\partial f}{\partial p} \right)^\top \lambda, \quad (49)$$

the second of which is the quadrature of (25) with $q(T) = 0$. CVODES carries it alongside the backward state, and the generated glue places it inside the error test rather than leaving it uncontrolled. Both equations are integrated from T down to t_0 , and the minus signs of (49) cancel against the reversed limits, so $q(t_0)$ is the parameter part of the gradient and $\lambda(t_0)$ the state part. Both right-hand sides are contractions of quantities the forward path already builds; only the index order differs.

Checkpoints and the forward state. The backward right-hand side needs $x(t)$, which the forward pass no longer holds. CVODES stores a checkpoint every `asaCheckpoints` accepted steps, 200 by default, and recovers the states between two checkpoints by re-integrating forward from the earlier one and interpolating. The argument trades memory against that repeated re-integration. This is also the substantive difference between the two discretisations: the states the adjoint sees are interpolated rather than the ones the forward run stepped through, so the resulting gradient is the derivative of no discretisation that was actually carried out. It approximates the exact gradient to $O(\text{tol})$ and differs by the same order from the derivative of the value the call reported.

The seed enters as a sequence of jumps. λ is initialised with the seed row of the last observation, and the sweep stops at every earlier observation time to add that row, which is the jump condition in (22). Since a jump constitutes a new initial condition, the backward problem is re-initialised wherever a row is nonzero. Re-initialising at every output time would instead reduce the method to first order at times the objective does not observe. The backward problem is a BDF solve with a dense linear solver, run at the caller's tolerances and with the caller's step budget rather than the CVODES default.

One backward solve per cotangent. Several seed columns reuse the same backward problem, but each requires a separate backward integration: the checkpoints are shared, the trajectory is not. The native mode sweeps its stored checkpoints once per column instead, and its `keepStore`, `store`, `adjointGrid` and `errWeights` arguments have no counterpart here, since CVODES holds its own checkpoints and runs its own backward controller.

Restrictions. Events and root functions are rejected at code-generation time rather than approximated, since a jump in the state implies a jump in λ and ASA provides no point at which the backward solver can be informed of one. Forward sensitivities cannot be requested for the same model, and there is no second-order mode. An event model, a Hessian, or a gradient consistent with the value it accompanies therefore requires the native backend.

Role. ASA is an independent implementation of the same mathematics, which makes it a check on the in-tree adjoint: a systematic error in either would appear as a disagreement between them, and the two agree to 10^{-6} on the test models. As a production route it is the slower of the two, by a factor of 1.9 on the chain of Table 4, where it remains a factor of five below the forward sensitivities it replaces.

4 cppFUN()

The `cppFUN()` entry point compiles a system of algebraic expressions

$$y_i = g_i(x_1, \dots, x_{n_x}, p_1, \dots, p_{n_p}), \quad i = 1, \dots, n_y, \quad (50)$$

together with optional first- and second-order derivatives. No time integration is performed: each call evaluates g at one or more rows of (x, p) values and returns the corresponding outputs, Jacobians, and Hessians as requested. Two use cases motivate the function. The first is the evaluation of observation maps $y = g(x, p)$ that connect integrator output to data in likelihood-based inference; the mapping is applied at every observation point and its Jacobian enters the gradient of the log-likelihood. The second is the evaluation of reparametrisation Jacobians $\partial p / \partial \theta$ that supply the seed matrix in equation (12) when the dynamic model is fitted in a transformed parameter space.

4.1 Unified API and two computational paths

Both values of the `derivMode` argument expose the same R-side methods: `func`, `jac`, `hess`, and the combined `evaluate`. They differ only in how the underlying C entry points are generated and dispatched.

- "symbolic": the analytic Jacobian and Hessian of g are derived by SymPy at code-generation time. The compiled module exposes `_eval`, `_jacobian` (when `deriv`), and `_hessian` (when `deriv2`). When seeds are supplied, the chain rule is applied by BLAS-3 DGEMM through two extern "C" wrappers `_chain_jac` and `_chain_hess`, which delegate to the inline routines in `cppde_chain_blas.hpp`.

- "forward" (default): no SymPy step. The compiled module exposes `_eval`, `_eval_ad` (when `deriv`, on `cppde::dual<double, 0>`), and `_eval_ad2` (when `deriv2`, on the nested-dual alias `cppde::dual2nd<double, 0>`). The chain rule is applied inside the AD evaluation by seeding the dual-number tangent slots; the same source that evaluates g over $\mathbb{R}$ then also evaluates the chain-ruled $\partial y / \partial \theta$ and $\partial^2 y / \partial \theta^2$.
- "reverse" (default alongside "forward"): the vector-Jacobian product `_vjp`, which contracts a cotangent of the outputs into cotangents of the variables and parameters at a cost independent of the upstream parameter count. It contracts the symbolic Jacobian, and is separately requestable: a caller that only ever multiplies by the Jacobian should not compile it.

"symbolic" is a backend for the forward Jacobian rather than a direction of its own, and cannot be combined with the other two; the vector-Jacobian product is symbolic either way. "forward" and "reverse" are independent build products and may be asked for together, which is the default.

The two paths are equivalent in exact arithmetic. The forward AD path avoids the SymPy expansion of $\partial g / \partial p$ and $\partial^2 g / \partial p^2$, whose intermediate expressions can grow sharply with the depth of the compositions in g . The symbolic path benefits from SymPy's common-subexpression elimination when the analytic derivatives remain compact. The reverse entry is described in Section 2 for the integrator; here it is the same construction without the time stepping, and one sweep over a straight-line program.

Call	dual path	symbolic path
<code>func(x, p)</code>	<code>_eval</code>	<code>_eval</code>
<code>jac(x, p)</code> (no seeds)	<code>_eval_ad</code> with identity seed	<code>_jacobian</code>
<code>jac(x, p, dX, dP)</code>	<code>_eval_ad</code> with user seeds	<code>_jacobian + _chain_jac</code>
<code>hess(x, p)</code> (no seeds)	<code>_eval_ad2</code> with identity seed	<code>_hessian</code>
<code>hess(x, p, dX, dP[, dX2, dP2])</code>	<code>_eval_ad2</code> with all seeds	<code>_hessian + _chain_hess</code>
<code>evaluate(..., deriv2 = FALSE)</code>	one <code>_eval_ad</code> -pass	<code>_eval + _jacobian (+ chain)</code>
<code>evaluate(..., deriv2 = TRUE)</code>	one <code>_eval_ad2</code> -pass	<code>_eval + _jacobian + _hessian (+ chain)</code>

Table 6: Dispatch within `cppFUN()`. Both modes accept the same arguments and return arrays of the same shape; only the C entries called and the location of the chain-rule contraction differ.

The combined entry point `evaluate()` is provided because in the dual path a single nested-dual evaluation produces the value, the first-order sensitivities, and the second-order sensitivities together; obtaining all three through separate calls to `func`, `jac`, and `hess` would repeat the underlying evaluation three times. In the symbolic path `evaluate()` simply combines the three independent C calls and applies the chain rule once.

4.2 Chain rule via dX, dP, dX2, dP2

The seed arguments encode the Jacobian and (optionally) Hessian of an upstream parameter map $\Phi : \theta \mapsto (x, p)$:

- `dX[obs, name, theta]` – $\partial x_i / \partial \theta_a$, indexed per observation row, per state name, per direction θ_a .
- `dP[name, theta]` – $\partial p_j / \partial \theta_a$, constant across observations.
- `dX2[obs, name, theta1, theta2]` – second-order sensitivities $\partial^2 x_i / \partial \theta_a \partial \theta_b$, optional. Defaults to zero (linear / affine Φ).
- `dP2[name, theta1, theta2]` – second-order parameter sensitivities $\partial^2 p_j / \partial \theta_a \partial \theta_b$, optional. Defaults to zero.

Symbols listed in `fixed` have their seed rows zeroed before the contraction; symbols absent from the seeds also contribute zero, not an error.

In the symbolic path the seeds are bundled into one matrix $S \in \mathbb{R}^{n_{\text{obs}} \times n_{\text{diff}} \times n_{\theta}}$ over the canonical-symbol basis

$n_{\text{diff}} = n_x + n_p$. The chain rule for the Jacobian is

$$\left. \frac{\partial y}{\partial \theta} \right|_{\text{obs}} = J(x_{\text{obs}}, p) \cdot S_{\text{obs}}, \quad (51)$$

implemented as one DGEMM per observation in `cppde::chain_jac`. For the Hessian, two contributions combine,

$$\left. \frac{\partial^2 y}{\partial \theta \partial \theta^\top} \right|_{\text{obs}, o} = S_{\text{obs}}^\top H_o(x_{\text{obs}}, p) S_{\text{obs}} + \sum_i J_{o,i}(x_{\text{obs}}, p) S2_{\text{obs}, i, \cdot, \cdot}, \quad (52)$$

implemented as two DGEMMs plus one DGEMV per (obs, output) in `cppde::chain_hess`. The second term vanishes when Φ is affine (`dX2` and `dP2` unset).

In the dual path the same contractions are obtained automatically by seeding the dual-number tangent slots. Each input scalar x_i becomes

$$\tilde{x}_i = (x_i + dX_{i,\cdot} \varepsilon_{\text{in}}) + (dX_{i,\cdot} + dX2_{i,\cdot,\cdot} \varepsilon_{\text{in}}) \varepsilon_{\text{out}},$$

i.e. a nested dual whose inner tangent slots carry $\partial x_i / \partial \theta$ and whose outer tangent slot k is itself an inner dual carrying $\partial x_i / \partial \theta_k$ as its value and $\partial^2 x_i / \partial \theta_k \partial \theta$ as its inner tangents. The dual-number multiplication rule then propagates both the $S^\top H S$ and the $J \cdot S2$ terms in one nested-arithmetic evaluation; the inner-tangent values of $\tilde{y}$ are read out as $\partial y / \partial \theta$, and the cross slots $\tilde{y}[k][m]$ as $\partial^2 y / \partial \theta_k \partial \theta_m$.

4.3 Pass-through of unmodelled inputs

The `attach.input` flag controls whether unmodelled columns of the input are appended to the output. When `attach.input = TRUE`, columns of the input matrix that do not appear among the named variables of (50) are passed through to the output unchanged: their value rows are copied, their Jacobian rows are zero for variables and the identity for parameters, and their Hessian rows are zero. The convention removes the need to extend g with no-op identities when the observation map does not consume every column of the integrator output, and matches the row-name-driven partial-shape `sens1ini` convention used by `solveODE()`.

5 Benchmarks

Both backends solve the same problem from the same symbolic input. Right-hand side, analytic Jacobian, events, linear-solver selection, and tolerances are shared, so a head-to-head measurement isolates the numerical core rather than the model formulation.

Setup. The suite comprises 41 problems: 31 published systems-biology models from a PETab benchmark collection and 10 classical test problems (van der Pol at $\mu = 1000$, Robertson, OREGO, HIRES, E5, Pollution, and method-of-lines semidiscretisations of the one-dimensional Brusselator and FitzHugh–Nagumo systems at $N = 24$ and $N = 64$). The state dimension ranges from $n_x = 2$ to 128, the number of differentiated parameters from $M = 1$ to 64. One further model of the collection, Lang_PLOSCoMputBiol2024 at $n_x = 124$, is left out: its solves alone take longer than the rest of the suite together, and it belongs to none of the regimes the comparison turns on. Both backends run their BDF/NDF method with the automatically selected linear solver at $(\text{rtol}, \text{atol}) \in \{(10^{-4}, 10^{-6}), (10^{-7}, 10^{-9}), (10^{-10}, 10^{-12})\}$, without and with first-order sensitivities. Each cell repeats the solve until one batch exceeds a wall-clock floor and reports the median of five such batches, on an otherwise idle single core. Figure 1 shows the resulting ratio $r = t_{\text{cvsode}}/t_{\text{cppode}}$ per problem, aggregated over the three tolerances by geometric mean.




Figure 1: Head-to-head ratio of the two backends, one marker per problem and mode; above the line `cppODE()` is faster. The abscissa is the number of parameters the model differentiates, which the plain solve does not use. Its markers are placed at the same abscissa as their sensitivity counterparts, and their flat course shows that the M dependence belongs to the sensitivity mode alone.

Without sensitivities the in-tree backend is faster on 38 of the 41 problems, by a geometric mean of 1.20, and the ratio is flat in M . The two backends take the same steps: the geometric means of the accepted-step and right-hand-side evaluation ratios are 1.01 and 1.05, which rules out a margin bought by integrating more loosely. This is expected, since the controller, the order selection, and the error constants of the in-tree multistep are those of CVODE, and the margin is therefore a cost-per-step effect. It arises from model specialisation: the right-hand side and the Jacobian are emitted as straight-line C++ and instantiated into the stepper as a functor template, so the calls inline, whereas CVODES reaches them through a function pointer with a `void*` user-data argument; the in-tree state vectors are contiguous with the extent known at the call site, whereas every SUNDIALS vector operation is dispatched through the operation table of the `N_Vector` it acts on. Part of the margin belongs to the code generator rather than to the integrator: common-subexpression elimination is applied to the in-tree right-hand side and Jacobian but not to the CVODES glue layer.

With sensitivities the geometric mean rises to 1.84, and Figure 1 shows that it does so as a function of M : from 1.17 over the nine problems with $M \leq 3$ to 3.26 over the nine with $13 \leq M \leq 24$. CVODES is driven in staggered mode with an analytic sensitivity right-hand side that is invoked once per differentiated parameter, each invocation forming one column of (12) and with it every structurally nonzero entry of J ; since the M invocations are separate calls, those entries are recomputed M times per corrector iteration. The in-tree backend obtains the same M columns from a single evaluation of f on a dual-typed state carrying all M tangents, and forms the shared subexpressions of J once. It reaches the same trajectories with roughly one twelfth of the right-hand-side evaluations.

Where the ordering reverses. The dual evaluation carries a cost per state component that is independent of M : the arithmetic on the value part, the tangent-slab bookkeeping, and the loss of the vectorisation the scalar loop admits. The in-tree sensitivity work therefore scales as $n_x(a + bM)$ with $a > 0$, whereas the CVODES work scales as $M \cdot \text{nnz}(J)$. For a method-of-lines discretisation $\text{nnz}(J) = \mathcal{O}(n_x)$, the state dimension cancels, and the comparison reduces to $a + bM$ against $\mathcal{O}(M)$: small M favours CVODES, and the more so the larger n_x is, since a is paid per component and amortised only over M . Every problem that falls below the parity line of Figure 1 with sensitivities is of this shape, and all of them lie at $M \leq 6$: SalazarCavazos_MBoC2020 ($n_x = 75$, $M = 6$) at 0.64 and the Brusselator at $N = 64$ and $N = 24$ ($M = 3$) at 0.67 and 0.81. Three problems fall below it without sensitivities, two of them by less than 10%.

Linear algebra and accuracy. Fifteen of the models were additionally solved with the linear solver pinned to dense and to sparse. The gain of the sparse path grows with the state dimension in the same way for both backends, from 1.08 at $n_x = 8$ to 2.78 at $n_x = 128$, and the auto-detection selected the faster path in every swept case, so Figure 1 is a comparison at matched linear algebra. Every solve was additionally scored against a reference trajectory at $\text{atol} = 10^{-14}$, $\text{rtol} = 10^{-12}$. Per problem, the deviation of the in-tree backend is a median 1.06 times that of CVODE, with the central 90% of the problems between 0.51 and 1.35. The two are not on equal footing there: the reference is itself a CVODE solution, so CVODE is measured against a tighter run of the same method through the same code, whose truncation error is structurally its own, while the in-tree backend is measured across implementations. The bias runs in CVODE’s favour, which makes the quoted ratio an upper bound. Four of the 852 cells failed and are excluded from the aggregates, all of them the in-tree backend with sensitivities at $\text{rtol} = 10^{-10}$: Oliveira_NatCommun2021 on each of its three linear-solver settings, and Weber_BMC2015. Repeating the suite on eight workers of the same machine, which inflates the absolute times through shared cache and reduced clocks but leaves the ratios intact, reproduces the two aggregates as 1.22 and 1.91, and the per-problem values with a correlation of 0.96 in $\log r$ with sensitivities. Without sensitivities the correlation drops to 0.52, since a substantial fraction of the plain solves complete in well under a millisecond and the measurement is then dominated by the contention the extra workers create. The problems that disagree beyond that scatter are all of that kind, HIRES at 1.42 against 0.77 the widest among them.

References

- [1] A. C. Hindmarsh *et al.*, “SUNDIALS: Suite of nonlinear and differential/algebraic equation solvers,” *ACM Transactions on Mathematical Software (TOMS)*, vol. 31, no. 3, pp. 363–396, 2005, doi: 10.1145/1089014.1089020.
- [2] L. F. Shampine and M. W. Reichelt, “The MATLAB ODE suite,” *SIAM Journal on Scientific Computing*, vol. 18, no. 1, pp. 1–22, 1997, doi: 10.1137/S1064827594276424.
- [3] E. Hairer and G. Wanner, *Solving ordinary differential equations II: Stiff and differential-algebraic problems*. in Springer series in computational mathematics. Springer Berlin Heidelberg, 2010. Available: <https://books.google.de/books?id=zZSv3acps1QC>
- [4] Ch. Tsitouras, “Runge–Kutta pairs of order 5(4) satisfying only the first column simplifying assumption,” *Computers & Mathematics with Applications*, vol. 62, no. 2, pp. 770–775, 2011, doi: 10.1016/j.camwa.2011.06.002.
- [5] E. Hairer, S. P. Nørsett, and G. Wanner, *Solving ordinary differential equations I: Nonstiff problems*. in Springer series in computational mathematics. Springer Berlin Heidelberg, 2008. Available: <https://books.google.de/books?id=F93u7VcSRyYC>
- [6] Y. Cao, S. Li, L. Petzold, and R. Serban, “Adjoint sensitivity analysis for differential-algebraic equations: The adjoint DAE system and its numerical solution,” *SIAM Journal on Scientific Computing*, vol. 24, no. 3, pp. 1076–1089, 2003, doi: 10.1137/S1064827501380630.

A BDF and NDF coefficients

Source. [2], Table 1; cppDE implementation in `cppde_multistepper.hpp`.

The Numerical Differentiation Formula (NDF) of order q modifies the classical BDF q corrector by adding a term proportional to the Nordsieck predictor error,

$$\sum_{m=1}^q \frac{1}{m} \nabla^m x_n - h_n f(x_n, t_n) - \kappa_q \gamma_q (x_n - x_n^{\text{pred}}) = 0, \quad \gamma_q = \sum_{m=1}^q \frac{1}{m}.$$

Setting $\kappa_q = 0$ recovers BDF q . The values κ_q adopted by cppDE (selectable at compile time via the argument `useNDF`, default `TRUE`) are tabulated in Table 7, together with the empirical step-size gain over classical BDF q and the resulting stability angle.

q	κ_q	Step-size gain over BDF q	Stability angle
1	-0.1850	+26 %	A-stable
2	-0.1111	+26 %	A-stable
3	-0.0823	+26 %	80°
4	-0.0415	+12 %	66°
5	0.0000	0 %	51° (= BDF5)

Table 7: NDF correction coefficients κ_q adopted by cppDE, the empirical step-size gain over classical BDF q at matched local error, and the resulting stability angle.

NDF5 is set to $\kappa_5 = 0$ because the BDF5 stability angle of 51° leaves no headroom for further reduction; NDF5 is therefore identical to BDF5.

B Adams-Moulton coefficients

Source. [1], ℓ -vector formulation as in CVODE; cppDE implementation in `adams_set_coefficients()` inside `cppde_multistepper.hpp`.

In contrast to BDF, the Adams-Moulton coefficients depend on the recent step-size history and are recomputed every step from the Nordsieck array. Define $\xi_j = (t_n - t_{n-j})/h_n$; the ℓ -vector recurrence reads

$$\ell_0(q) = 1, \quad \ell_j(q) = \ell_{j-1}(q) + \xi_j^{-1}, \quad j = 1, \dots, q,$$

followed by a polynomial expansion that yields the corrector weights $\beta_0, \dots, \beta_q$ used in the current step. Because the coefficients are step-size adaptive, no static table can reproduce them in general; the canonical fixed-step values for $h_n = h_{n-1} = \dots$ are given in standard references [5], Table III.1.1 and reduce to the recurrence above in that limit.

The maximum order in cppDE is $q_{\max} = 12$, matching CVODE's default; the order controller may select any $q \in \{1, \dots, 12\}$ at run time.

C Rosenbrock4 coefficients

Source. [3, Sec. IV.7]; identical to the default coefficients in the `rosenbrock4` stepper of Boost.Numeric.Odeint. cppDE implementation in `default_rosenbrock_coefficients` in `cppde_rosenbrock4.hpp`.

A Rosenbrock method of s stages with diagonal coefficient γ is defined, with iteration matrix $W = (\gamma h_n)^{-1} I -$

$J(x_n, t_n)$, by the stage equations

$$W k_i = f\left(x_n + \sum_{j<i} \alpha_{ij} k_j, t_n + c_i h_n\right) + \frac{1}{h_n} \sum_{j<i} \frac{\gamma_{ij}}{\gamma} k_j + h_n d_i \frac{\partial f}{\partial t},$$

$$x_{n+1} = x_n + \sum_{i=1}^s m_i k_i, \quad \hat{x}_{n+1} = x_n + \sum_{i=1}^s \hat{m}_i k_i.$$

The cppDE parameterisation uses $s = 6$ stages with $\gamma = 1/4$. The numerical values of all coefficients are reproduced in the following tables, indexed exactly as in the source code: the diagonal coefficient and the time partials and node positions in Table 8, the stage couplings α_{ij} and γ_{ij} in Tables 9 and 10, the sixth-row weights in Table 11, and the dense-output coefficients in Table 12.

C.1 Diagonal coefficient, time partials, and node positions

γ	0.25	d_1	0.25	c_2	0.386
d_2	-0.1043	d_3	0.1035	c_3	0.21
d_4	0.0362			c_4	0.63

Table 8: Rosenbrock4 diagonal coefficient γ , time-partial coefficients d_i , and stage node positions c_i .

C.2 Stage couplings α_{ij}

	$j = 1$	$j = 2$	$j = 3$	$j = 4$
$\alpha_{2,j}$	1.5440000000			
$\alpha_{3,j}$	0.9466785281	0.2557011699		
$\alpha_{4,j}$	3.3148251871	2.8961240160	0.9986419140	
$\alpha_{5,j}$	1.2212245092	6.0191344813	12.5370833293	-0.6878860361

Table 9: Rosenbrock4 stage-coupling coefficients α_{ij} . Row index i is the stage being formed; column index j ranges over earlier stages.

C.3 Stage couplings γ_{ij}

	$j = 1$	$j = 2$	$j = 3$	$j = 4$
$\gamma_{2,j}$	-5.6688000000			
$\gamma_{3,j}$	-2.4300933568	-0.2063599157		
$\gamma_{4,j}$	-0.1073529058	-9.5945622510	-20.4702861481	
$\gamma_{5,j}$	7.4964433140	-10.2468043146	-33.9999035282	11.7089089321

Table 10: Rosenbrock4 stage-coupling coefficients γ_{ij} , appearing in the $(1/h_n) \sum_{j<i} (\gamma_{ij}/\gamma) k_j$ term of the stage equation. Row and column conventions match Table 9.

C.4 Solution and error weights

The advanced solution and the embedded error estimate are formed from the stage values $k_1, \dots, k_5$ via

$$x_{n+1} = x_n + \sum_{i=1}^5 \alpha_{6,i} k_i, \quad e_{n+1} = \sum_{i=1}^5 \gamma_{6,i} k_i,$$

with the sixth-row coefficients listed in Table 11.

	$i = 1$	$i = 2$	$i = 3$	$i = 4$	$i = 5$
$\alpha_{6,i}$	1.2212245092	6.0191344813	12.5370833293	-0.6878860361	1.0
$\gamma_{6,i}$	8.0832467959	-7.9811329881	-31.5215943287	16.3193054312	-6.0588182388

Table 11: Rosenbrock4 sixth-row coefficients. The $\alpha_{6,i}$ values combine the stage values into the advanced solution x_{n+1} ; the $\gamma_{6,i}$ values supply the embedded order-3 error estimate e_{n+1} .

C.5 Dense-output coefficients

Continuous output on $[t_n, t_{n+1}]$ is constructed from the five intermediate stage values together with two auxiliary coefficient sets $d_i^{(2)}$ and $d_i^{(3)}$, listed in Table 12, that combine the stages with the step solution to produce an order-3 interpolant.

	$i = 1$	$i = 2$	$i = 3$	$i = 4$	$i = 5$
$d_i^{(2)}$	10.1262350834	-7.4879958776	-34.8009186156	-7.9927717076	1.0251377233
$d_i^{(3)}$	-0.6762803393	6.0877146517	16.4308432089	24.7672251142	-6.5943891257

Table 12: Rosenbrock4 dense-output coefficients $d_i^{(2)}$ and $d_i^{(3)}$. Together with the stage values they form the order-3 interpolant on $[t_n, t_{n+1}]$ used to serve user output points falling between accepted steps.

D Tsitouras 5(4) coefficients

Source. [4], Table 1 and equations (7)–(8); cppDE implementation in `cppde_tsit5.hpp`.

The Butcher tableau of the Tsitouras 5(4) pair, reproduced in Table 13, has $s = 7$ stages and obeys the simplifying assumption $a_{i,1} = 0$ for $i \geq 2$; the seventh stage is the FSAL stage. Both rows of weights are listed beneath the A matrix: b_i for the order-5 solution and $\hat{b}_i$ for the embedded order-4 estimate.

0							
0.161	0.161						
0.327	-0.0084806555	0.3354806555					
0.9	2.8971530571	-6.3594484900	4.3622954329				
0.9800255409	5.3258648284	-11.7488835641	7.4955393429	-0.0924950664			
1.0	5.8614554429	-12.9209693178	8.1593678986	-0.0715849733	-0.0282690504		
1.0	0.0964607668	0.01	0.4798896504	1.3790085741	-3.2900695154	2.3247105241	
b_i	0.0964607668	0.01	0.4798896504	1.3790085741	-3.2900695154	2.3247105241	0
$\hat{b}_i$	0.0946807557	0.0091835655	0.4877705284	1.2342975669	-2.7077123499	1.8666284182	0.0151515152

Table 13: Butcher tableau of the Tsitouras 5(4) pair. The first column lists the node positions c_i ; the centre block is the strictly lower-triangular coefficient matrix a_{ij} . The b_i row gives the order-5 weights, and the $\hat{b}_i$ row gives the embedded order-4 weights.

The seventh stage is $k_7 = f(x_{n+1}, t_n + h_n)$, which is re-used as k_1 at step $n + 1$. Because $b_7 = 0$, the order-5 solution does not depend on k_7 ; because $\hat{b}_7 \neq 0$, the embedded order-4 estimate does. The local-error estimator used by the step controller is

$$e_n = h_n \sum_{i=1}^7 (\hat{b}_i - b_i) k_i.$$

The seven differences $\hat{b}_i - b_i$ are stored in `cppde_tsit5.hpp` as the coefficients `e1`, ..., `e7` and reproduced in Table 14.

i	$\hat{b}_i - b_i$
1	-0.0017800111
2	-0.0008164345
3	0.0078808780
4	-0.1447110072
5	0.5823571655
6	-0.4580821059
7	0.0151515152

Table 14: Tsit5 embedded error coefficients $\hat{b}_i - b_i$, stored in `cppde_tsit5.hpp` as `e1`, ..., `e7`. The seven values sum to numerical zero, the consistency condition for the embedded pair.

The seven values sum to numerical zero, as required for the consistency of the embedded pair, and the row $a_{7,j}$ of the A matrix coincides with b_j for $j = 1, \dots, 6$, which is the FSAL property.